

Assignment2

October 21, 2020

1 Preprocessing of Discrete Data 1 - Cancer Dataset

```
[1]: #importing libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
# visualization
import seaborn as sns

from math import sqrt
```

```
[2]: # reading data
data = pd.read_csv('dataDiscrete/data.csv')
```

```
[3]: # first 5 entries of dataframe
data.head()
```

```
[3]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave	points_mean	\
0	0.11840	0.27760	0.3001		0.14710	

1	0.08474	0.07864	0.0869	0.07017
2	0.10960	0.15990	0.1974	0.12790
3	0.14250	0.28390	0.2414	0.10520
4	0.10030	0.13280	0.1980	0.10430

	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	17.33	184.60	2019.0	0.1622	
1	23.41	158.80	1956.0	0.1238	
2	25.53	152.50	1709.0	0.1444	
3	26.50	98.87	567.7	0.2098	
4	16.67	152.20	1575.0	0.1374	

	compactness_worst	concavity_worst	concave points_worst	symmetry_worst	\
0	0.6656	0.7119	0.2654	0.4601	
1	0.1866	0.2416	0.1860	0.2750	
2	0.4245	0.4504	0.2430	0.3613	
3	0.8663	0.6869	0.2575	0.6638	
4	0.2050	0.4000	0.1625	0.2364	

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN

[5 rows x 33 columns]

[]:

```
[4]: #data preprocessing
data_count=data.diagnosis.value_counts(normalize = True)
data_count = pd.Series(data_count)
data_count = pd.DataFrame(data_count)
data_count.index = ['Benign', 'Malignant']

data_count['Percent'] = 100*data_count['diagnosis']/sum(data_count['diagnosis'])
data_count['Percent'] = data_count['Percent'].round().astype('int')
data_count
```

```
[4]:      diagnosis  Percent
Benign    0.627417      63
Malignant  0.372583      37
```

```
[5]: # changing categorical data to numerical
data['diagnosis']= data['diagnosis'].map({'M':1, 'B':0})
```

```

# checking the different values contained in the diagnosis column
#Benign : 0
#Malign : 1

data['diagnosis'].value_counts()

# it shows that all the data are unique that is.
data['id'].nunique()

# here data in last column is empty and id is unique, so removing this does not
→affect data

data.drop(data.columns[[-1, 0]], axis=1, inplace=True)

data.shape

```

[5]: (569, 31)

```

[6]: ### Applying Dimensionality Reduction

from sklearn.preprocessing import StandardScaler # Import
scaler = StandardScaler() # Instantiate
scaler.fit(data) # Fit
scaled_data = scaler.transform(data) # Transform

# Applying PCA

from sklearn.decomposition import PCA # Import

pca = PCA(n_components=2) # Instantiate

pca.fit(scaled_data) # Fit

X_pca = pca.transform(scaled_data) # Transform

print(" Data dimensions before reduction:",scaled_data.shape)
print(" Data dimensions after reduction:",X_pca.shape)

```

```

Data dimensions before reduction: (569, 31)
Data dimensions after reduction: (569, 2)

```

```

[7]: # Splitting the data set into train and test set
from sklearn.model_selection import train_test_split

features = data.drop(columns = ['diagnosis'])

```

```

target = data['diagnosis']
X_train1, X_test1, y_train1, y_test1 = train_test_split(features, target,
    ↳test_size = 0.3, random_state = 0)

print ("Train data set size : ", X_train1.shape)
print ("Test data set size : ", X_test1.shape)

```

Train data set size : (398, 30)

Test data set size : (171, 30)

```

[8]: # applying 5NN , id3, Naive bayes
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
# visualization
import seaborn as sns
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import precision_score, recall_score
from math import sqrt
# Spot-Check Algorithms
models = []

models.append(( 'KNN' , KNeighborsClassifier(n_neighbors = 5, metric =
    ↳'minkowski', p = 2)))
models.append(( 'ID3' , DecisionTreeClassifier(criterion = 'entropy',
    ↳random_state = 0)))
models.append(( 'NB' , GaussianNB()))

# Test options and evaluation metric
num_folds = 5
num_instances = len(X_train1)
seed = 7
scoring = 'accuracy'

# Test options and evaluation metric
num_folds = 5
num_instances = len(X_train1)
seed = 7
scoring = 'accuracy'
results = []

```

```

names = []
print("Cancer dataset ")

for name, model in models:
    #applying stratified 5 fold technique
    cv_results = StratifiedKFold(n_splits=5, random_state=1)
    # calculating scores
    scores = cross_val_score(model, X_train1, y_train1, scoring='accuracy',
    ↪cv=cv_results, n_jobs=-1)
    model.fit(X_test1, y_test1)
    pred = model.predict(X_test1)

    results.append(model.score(X_test1, y_test1))
    names.append(name)
    msg = "%s: %f (%f)" % (name,scores.mean(), scores.std())
    print('\n')
#mean and standard deviation of prediction
    print('      Mean      Standard Deviation')
    print(msg)
#printing values for comparision ;ike confusion matrix and accuracy
    print('confusion matrix')
    print(confusion_matrix(y_test1,pred))

    # Print out confusion matrix
    cmat = confusion_matrix(y_test1, pred)
#print(cmat)
    print('TP - True Negative {}'.format(cmat[0,0]))
    print('FP - False Positive {}'.format(cmat[0,1]))
    print('FN - False Negative {}'.format(cmat[1,0]))
    print('TP - True Positive {}'.format(cmat[1,1]))
    Accuracy = (np.divide(np.sum([cmat[0,0],cmat[1,1]]),np.sum(cmat)))
    print('Accuracy Rate:',Accuracy)
    #interval = z * sqrt( (accuracy * (1 - accuracy)) / n)
    #confidence interval with 80%
    confidence_interval = 1.282 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
    print("True error (with 80% Confidence Interval): ", confidence_interval)
    confidence_interval2 = 1.645 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
    print("True error (with 90% Confidence Interval): ", confidence_interval2)
    # printing pessimistic error of the
    print('Pessimistic error: {}'.format(np.divide(np.
    ↪sum([cmat[0,1],cmat[1,0]]),np.sum(cmat))))
    print('\n')

print('→ 5-Fold cross-validation accuray score for the training data for_
    ↪above classifiers')
print('test results',results)

```

Cancer dataset

```
      Mean      Standard Deviation
KNN: 0.919679 (0.018417)
confusion matrix
[[106   2]
 [  4  59]]
TP - True Negative 106
FP - False Positive  2
FN - False Negative  4
TP - True Positive  59
Accuracy Rate: 0.9649122807017544
True error (with 80% Confidence Interval): 0.011824096788044459
True error (with 90% Confidence Interval): 0.015172105472958762
Pessimistic error: 0.03508771929824561
```

```
      Mean      Standard Deviation
ID3: 0.916923 (0.032763)
confusion matrix
[[108   0]
 [  0  63]]
TP - True Negative 108
FP - False Positive  0
FN - False Negative  0
TP - True Positive  63
Accuracy Rate: 1.0
True error (with 80% Confidence Interval): 0.0
True error (with 90% Confidence Interval): 0.0
Pessimistic error: 0.0
```

```
      Mean      Standard Deviation
NB: 0.942244 (0.014880)
confusion matrix
[[104   4]
 [  6  57]]
TP - True Negative 104
FP - False Positive  4
FN - False Negative  6
TP - True Positive  57
Accuracy Rate: 0.9415204678362573
True error (with 80% Confidence Interval): 0.015078679724556699
```

True error (with 90% Confidence Interval): 0.019348227883694048
Pessimistic error: 0.05847953216374269

-> 5-Fold cross-validation accuracy score for the training data for above classifiers

test results [0.9649122807017544, 1.0, 0.9415204678362573]

```
[9]: # applying ripper algorithm
import wittgenstein as lw
ripper_clf = lw.RIPPER() # Or irep_clf = lw.IREP() to build a model using IREP
ripper_clf.fit( X_train1, y_train1) # Or pass X and y data to .fit
ripper_clf
```

```
[9]: <RIPPER(dl_allowance=64, k=2, max_rules=None, prune_size=0.33,
random_state=None, verbosity=0, max_rule_conds=None, max_total_conds=None,
n_discretize_bins=10)>
```

```
[10]: # ruleset of ripper algo
ripper_clf.ruleset_
```

```
[10]: <Ruleset [radius_mean=17.93-20.26] V [perimeter_worst=119.4-140.9] V
[radius_mean=20.26-28.11] V [concavity_mean=0.15-0.22] V
[perimeter_worst=105.9-119.4^radius_mean=13.46-14.25] V
[texture_worst=30.96-34.85^radius_mean=15.34-17.93] V
[perimeter_worst=105.9-119.4^texture_worst=28.74-30.96] V
[texture_mean=21.38-22.55^radius_mean=14.25-15.34] V
[compactness_mean=0.19-0.31]>
```

```
[11]: pred = ripper_clf.predict(X_test1)

print("Confusion Matrix")
print(confusion_matrix(y_test1,pred))
print('\n')
print("classification report")
print(classification_report(y_test1,pred))
# Print out confusion matrix
cmat = confusion_matrix(y_test1, pred)
#print(cmat)
print('TP - True Negative {}'.format(cmat[0,0]))
print('FP - False Positive {}'.format(cmat[0,1]))
print('FN - False Negative {}'.format(cmat[1,0]))
print('TP - True Positive {}'.format(cmat[1,1]))
Accuracy = (np.divide(np.sum([cmat[0,0],cmat[1,1]]),np.sum(cmat)))
print('Accuracy Rate:',Accuracy)
#interval = z * sqrt( (accuracy * (1 - accuracy)) / n)
#confidence interval with 80%
```

```

confidence_interval = 1.282 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 80% Confidence Interval): ", confidence_interval)
confidence_interval2 = 1.645 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 90% Confidence Interval): ", confidence_interval2)
print('Pessimistic Error: {}'.format(np.divide(np.sum([cmat[0,1],cmat[1,0]]),np.
↪sum(cmat))))
print('\n')

```

Confusion Matrix

```

[[102   6]
 [  9  54]]

```

classification report

	precision	recall	f1-score	support
0	0.92	0.94	0.93	108
1	0.90	0.86	0.88	63
accuracy			0.91	171
macro avg	0.91	0.90	0.90	171
weighted avg	0.91	0.91	0.91	171

TP - True Negative 102

FP - False Positive 6

FN - False Negative 9

TP - True Positive 54

Accuracy Rate: 0.9122807017543859

True error (with 80% Confidence Interval): 0.018178511005578008

True error (with 90% Confidence Interval): 0.02332578050247724

Pessimistic Error: 0.08771929824561403

2 Preprocessing of Continuous Data - Census/Adult Income dataset

```

[12]: # Import libraries necessary for this project
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import seaborn as sns
sns.set(style="darkgrid")
from time import time

```



```

from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import OneHotEncoder
import scipy.stats as stats
from scipy.stats import kurtosistest

# displaying for notebooks
%matplotlib inline

columns = ['age', 'workclass', 'fnlwgt', 'education', 'education-num',
           'marital-status', 'occupation',
           'relationship', 'race', 'sex', 'capital-gain', 'capital-loss',
           'hours-per-week', 'native-country', 'income']

# Load the Census dataset
census = pd.read_csv('dataConti/adult.data', header=None, names=columns,
                    skipinitialspace=True)

# Success - Display the first 5 record
display(census.head())

```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

```

[13]: import warnings
warnings.filterwarnings("ignore")

# Drop the fnlwgt column which is useless for later analysis

```

```

census = census.drop('fnlwgt', axis=1)

# Read in test data
census_test = pd.read_csv('dataConti/adult.test', header=None, skiprows=1,
    ↪names=columns, skipinitialspace=True)

# Drop the fnlwgt column which is useless for later analysis
census_test = census_test.drop('fnlwgt', axis=1)

# Remove '.' in income column
census_test['income'] = census_test['income'].apply(lambda x: '>50K' if
    ↪x=='>50K.' else '<=50K')

# Review several rows and shape of data set
display(census_test.head())
display(census_test.shape)

```

	age	workclass	education	education-num	marital-status	\
0	25	Private	11th	7	Never-married	
1	38	Private	HS-grad	9	Married-civ-spouse	
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	
3	44	Private	Some-college	10	Married-civ-spouse	
4	18	?	Some-college	10	Never-married	

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	Machine-op-inspct	Own-child	Black	Male	0	0	
1	Farming-fishing	Husband	White	Male	0	0	
2	Protective-serv	Husband	White	Male	0	0	
3	Machine-op-inspct	Husband	Black	Male	7688	0	
4	?	Own-child	White	Female	0	0	

	hours-per-week	native-country	income
0	40	United-States	<=50K
1	50	United-States	<=50K
2	40	United-States	>50K
3	40	United-States	>50K
4	30	United-States	<=50K

(16281, 14)

```

[14]: # Convert '?' to NaNs and remove the entries with NaN value
object_col = census.select_dtypes(include=object).columns.tolist()
for col in object_col:
    census.loc[census[col]=='?', col] = np.nan
    census_test.loc[census_test[col]=='?', col] = np.nan

```

```
# Perform an mssing assessment in each column of the dataset.
col_missing_pct = census.isna().sum()/census.shape[0]
col_missing_pct.sort_values(ascending=False)
```

```
[14]: occupation      0.056601
      workclass       0.056386
      native-country  0.017905
      income          0.000000
      hours-per-week  0.000000
      capital-loss    0.000000
      capital-gain    0.000000
      sex            0.000000
      race           0.000000
      relationship    0.000000
      marital-status  0.000000
      education-num   0.000000
      education       0.000000
      age            0.000000
      dtype: float64
```

```
[15]: # Removing data entries with missing value
adult_train = census.dropna(axis=0, how='any')
adult_test = census_test.dropna(axis=0, how='any')

# Show the results of the split
print("After removing the missing value:")
print("Training set has {} samples.".format(adult_train.shape[0]))
print("Testing set has {} samples.".format(adult_test.shape[0]))
```

After removing the missing value:
Training set has 30162 samples.
Testing set has 15060 samples.

```
[16]: num_col = adult_train.dtypes[adult_train.dtypes != 'object'].index

# Split the data into features and target label
income_raw = adult_train['income']
feature_raw = adult_train.drop('income', axis=1)

income_raw_test = adult_test['income']
feature_raw_test = adult_test.drop('income', axis=1)

# Log transform the skewed feature highly-skewed feature 'capital-gain' and
↪ 'capital-loss'.
skewed = ['capital-gain', 'capital-loss']
census_log = pd.DataFrame(data=feature_raw)
census_log[skewed] = feature_raw[skewed].apply(lambda x: np.log(x + 1))
```

```

census_log_test = pd.DataFrame(data=feature_raw_test)
census_log_test[skewed] = feature_raw_test[skewed].apply(lambda x: np.log(x +
→1))

# Initialize a scaler, then apply it to the features
scaler = MinMaxScaler() # default=(0, 1)

features_log_minmax_transform = pd.DataFrame(data = census_log)
features_log_minmax_transform[num_col] = scaler.
→fit_transform(census_log[num_col])

# Transform the test data set
features_log_minmax_transform_test = pd.DataFrame(data = census_log_test)
features_log_minmax_transform_test[num_col] = scaler.
→transform(census_log_test[num_col])

# Show an example of a record with scaling applied
display(features_log_minmax_transform.head())
display(features_log_minmax_transform_test.head())

```

	age	workclass	education	education-num	marital-status \
0	0.301370	State-gov	Bachelors	0.800000	Never-married
1	0.452055	Self-emp-not-inc	Bachelors	0.800000	Married-civ-spouse
2	0.287671	Private	HS-grad	0.533333	Divorced
3	0.493151	Private	11th	0.400000	Married-civ-spouse
4	0.150685	Private	Bachelors	0.800000	Married-civ-spouse

	occupation	relationship	race	sex	capital-gain \
0	Adm-clerical	Not-in-family	White	Male	0.667492
1	Exec-managerial	Husband	White	Male	0.000000
2	Handlers-cleaners	Not-in-family	White	Male	0.000000
3	Handlers-cleaners	Husband	Black	Male	0.000000
4	Prof-specialty	Wife	Black	Female	0.000000

	capital-loss	hours-per-week	native-country
0	0.0	0.397959	United-States
1	0.0	0.122449	United-States
2	0.0	0.397959	United-States
3	0.0	0.397959	United-States
4	0.0	0.397959	Cuba

	age	workclass	education	education-num	marital-status \
0	0.109589	Private	11th	0.400000	Never-married
1	0.287671	Private	HS-grad	0.533333	Married-civ-spouse
2	0.150685	Local-gov	Assoc-acdm	0.733333	Married-civ-spouse

3	0.369863	Private	Some-college	0.600000	Married-civ-spouse
5	0.232877	Private	10th	0.333333	Never-married

	occupation	relationship	race	sex	capital-gain	capital-loss \
0	Machine-op-inspct	Own-child	Black	Male	0.000000	0.0
1	Farming-fishing	Husband	White	Male	0.000000	0.0
2	Protective-serv	Husband	White	Male	0.000000	0.0
3	Machine-op-inspct	Husband	Black	Male	0.777174	0.0
5	Other-service	Not-in-family	White	Male	0.000000	0.0

	hours-per-week	native-country
0	0.397959	United-States
1	0.500000	United-States
2	0.397959	United-States
3	0.397959	United-States
5	0.295918	United-States

```
[17]: X = feature_raw

y =income_raw

# splitting data to train and test for further preprocessing

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳random_state = 0)

from sklearn import preprocessing

categorical = ['workclass', 'education', 'marital-status', 'occupation',
↳'relationship', 'race', 'sex', 'native-country']
for feature in categorical:
    le = preprocessing.LabelEncoder()
    X_train[feature] = le.fit_transform(X_train[feature])
    X_test[feature] = le.transform(X_test[feature])

# Changing the income column into Numerical Value
y_train = y_train.map({'<=50K':0, '>50K':1})
```

```
[18]: # applying 5NN , id3, Naive bayes
```

```
# Spot-Check Algorithms
models = []
```

```

models.append(( 'KNN' , KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)))
models.append(( 'ID3' , DecisionTreeClassifier(criterion = 'entropy', random_state = 0)))
models.append(( 'NB' , GaussianNB()))

# Test options and evaluation metric
num_folds = 5
num_instances = len(X_train.dtypes[X_train.dtypes != 'object'].index)
seed = 7
scoring = 'accuracy'

# Test options and evaluation metric
num_folds = 5
num_instances = len(X_train)
seed = 7
scoring = 'accuracy'
results = []
names = []
print("For Census dataset ")

for name, model in models:

    cv_results = StratifiedKFold(n_splits=5, random_state=1)
    scores = cross_val_score(model, X_train, y_train, scoring='accuracy', cv=cv_results, n_jobs=-1)

    model.fit(X_test, y_test)
    pred = model.predict(X_test)

    results.append(model.score(X_test, y_test))
    names.append(name)
    msg = "%s: %f (%f)" % (name, scores.mean(), scores.std())
    print('      Mean      Standard Deviation')
    print(msg)
    print('confusion matrix')
    print(confusion_matrix(y_test, pred))
    # Print out confusion matrix
    cmat = confusion_matrix(y_test, pred)
    #print(cmat)
    print('TP - True Negative {}'.format(cmat[0,0]))
    print('FP - False Positive {}'.format(cmat[0,1]))
    print('FN - False Negative {}'.format(cmat[1,0]))
    print('TP - True Positive {}'.format(cmat[1,1]))
    #Accuracy Rate

```

```

Accuracy = (np.divide(np.sum([cmat[0,0],cmat[1,1]]),np.sum(cmat)))
print('Accuracy Rate:',Accuracy)
#interval = z * sqrt( (accuracy * (1 - accuracy)) / n)
    #confidence interval with 80%
confidence_interval = 1.282 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 80% Confidence Interval): ", confidence_interval)
confidence_interval2 = 1.645 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 90% Confidence Interval): ", confidence_interval2)
    # printing pessimistic error of the
    #pessimistic error
print('Pessimistic Error: {}'.format(np.divide(np.
→sum([cmat[0,1],cmat[1,0]]),np.sum(cmat))))
print('\n')

print('→ 5-Fold cross-validation accuray score for the training data for_
→above classifiers')
print('test results',results)

```

For Census dataset

```

    Mean    Standard Deviation
KNN: 0.818216 (0.004867)
confusion matrix
[[6307  457]
 [ 792 1493]]
TP - True Negative 6307
FP - False Positive 457
FN - False Negative 792
TP - True Positive 1493
Accuracy Rate: 0.8619736987512432
True error (with 80% Confidence Interval):  0.0030432734954399443
True error (with 90% Confidence Interval):  0.003904980421215841
Pessimistic Error: 0.13802630124875676

```

```

    Mean    Standard Deviation
ID3: 0.810969 (0.009602)
confusion matrix
[[6756    8]
 [ 103 2182]]
TP - True Negative 6756
FP - False Positive 8
FN - False Negative 103
TP - True Positive 2182
Accuracy Rate: 0.9877334512100785
True error (with 80% Confidence Interval):  0.0009711674716237582
True error (with 90% Confidence Interval):  0.0012461548290336054
Pessimistic Error: 0.012266548789921538

```

```

        Mean      Standard Deviation
NB: 0.804481 (0.006848)
confusion matrix
[[5886  878]
 [ 935 1350]]
TP - True Negative 5886
FP - False Positive 878
FN - False Negative 935
TP - True Positive 1350
Accuracy Rate: 0.7996463697646149
True error (with 80% Confidence Interval):  0.003531511658043248
True error (with 90% Confidence Interval):  0.0045314638669899715
Pessimistic Error: 0.20035363023538513

```

-> 5-Fold cross-validation accuray score for the training data for above classifiers
test results [0.8619736987512432, 0.9877334512100785, 0.7996463697646149]

```

[19]: # applying ripper algorithm
import wittgenstein as lw
ripper_clf = lw.RIPPER() # Or irep_clf = lw.IREP() to build a model using IREP
ripper_clf.fit( X_train, y_train) # Or pass X and y data to .fit
ripper_clf

```

```

[19]: <RIPPER(dl_allowance=64, k=2, max_rules=None, prune_size=0.33,
random_state=None, verbosity=0, max_rule_conds=None, max_total_conds=None,
n_discretize_bins=10)>

```

```

[20]: # ruelset of ripper algo
ripper_clf.ruleset_

```

```

[20]: <Ruleset [marital-status=2^education-num=0.8-1.0^capital-gain=0.0-1.0^hours-per-
week=0.5-0.7] V [marital-status=2^education-num=0.6-0.8^capital-
gain=0.0-1.0^occupation=2.0-3.0] V [marital-status=2^education-
num=0.8-1.0^capital-gain=0.0-1.0^hours-per-week=0.4-0.5^occupation=7.0-9.0] V
[marital-status=2^education-num=0.8-1.0^capital-loss=0.0-0.98^hours-per-
week=0.4-0.5] V [marital-status=2^education-
num=0.6-0.8^occupation=2.0-3.0^hours-per-week=0.4-0.5^capital-loss=0.0-0.98] V
[marital-status=2^education-num=0.8-1.0^capital-
gain=0.0-1.0^education=11.0-15.0^hours-per-week=0.24-0.38] V [marital-
status=2^education-num=0.6-0.8^hours-per-
week=0.4-0.5^occupation=2.0-3.0^age=0.41-0.49] V [marital-status=2^education-
num=0.8-1.0^race=4^capital-loss=0.0-0.98] V [marital-status=2^education-
num=0.6-0.8^occupation=2.0-3.0^workclass=2^age=0.34-0.41^race=4^hours-per-

```


week=0.38-0.4] V [marital-status=2^education-num=0.6-0.8^hours-per-week=0.4-0.5^occupation=2.0-3.0] V [marital-status=2^education-num=0.6-0.8^capital-gain=0.0-1.0^hours-per-week=0.4-0.5^workclass=2^occupation=7.0-9.0] V [marital-status=2^education-num=0.8-1.0^occupation=2.0-3.0^hours-per-week=0.4-0.5] V [marital-status=2^occupation=7.0-9.0^education-num=0.8-1.0^hours-per-week=0.38-0.4] V [marital-status=2^education-num=0.6-0.8^capital-gain=0.0-1.0^occupation=9.0-11.0^age=0.49-0.63] V [marital-status=2^education-num=0.6-0.8^capital-gain=0.0-1.0^hours-per-week=0.5-0.7] V [marital-status=2^education-num=0.6-0.8^capital-loss=0.0-0.98] V [marital-status=2^education-num=0.6-0.8^capital-gain=0.0-1.0] V [marital-status=2^education=11.0-15.0^education-num=0.8-1.0^occupation=2.0-3.0^age=0.29-0.34] V [marital-status=2^occupation=7.0-9.0^education-num=0.8-1.0^workclass=4^education=11.0-15.0^hours-per-week=0.5-0.7] V [marital-status=2^education=11.0-15.0^education-num=0.8-1.0^occupation=2.0-3.0^relationship=0^workclass=2] V [marital-status=2^education-num=0.6-0.8^hours-per-week=0.4-0.5^workclass=3] V [marital-status=2^occupation=2.0-3.0^capital-gain=0.0-1.0^workclass=3^education=11.0-15.0] V [marital-status=2^occupation=7.0-9.0^education-num=0.8-1.0^workclass=3] V [marital-status=2^education-num=0.6-0.8^occupation=7.0-9.0^workclass=2^hours-per-week=0.4-0.5^age=0.29-0.34] V [marital-status=2^education-num=0.6-0.8^occupation=7.0-9.0^relationship=5^age=0.23-0.29] V [marital-status=2^occupation=2.0-3.0^capital-gain=0.0-1.0^age=0.41-0.49] V [marital-status=2^education-num=0.6-0.8^occupation=2.0-3.0^workclass=2^hours-per-week=0.5-0.7^age=0.18-0.23] V [marital-status=2^education=11.0-15.0^education-num=0.8-1.0^occupation=9.0-11.0^workclass=3] V [marital-status=2^education-num=0.6-0.8^occupation=7.0-9.0] V [marital-status=2^education=11.0-15.0^age=0.41-0.49^hours-per-week=0.4-0.5^workclass=2] V [marital-status=2^education=11.0-15.0^education-num=0.8-1.0^occupation=7.0-9.0^hours-per-week=0.4-0.5^age=0.29-0.34] V [marital-status=2^education=11.0-15.0^capital-gain=0.0-1.0^age=0.34-0.41] V [marital-status=2^occupation=2.0-3.0^capital-gain=0.0-1.0] V [marital-status=2^occupation=2.0-3.0^capital-loss=0.0-0.98^education=9.0-11.0] V [marital-status=2^education=11.0-15.0^education-num=0.8-1.0^occupation=2.0-3.0^workclass=1] V [marital-status=2^education=11.0-15.0^occupation=2.0-3.0^workclass=0] V [marital-status=2^education=11.0-15.0^education-num=0.8-1.0^occupation=7.0-9.0^race=4^relationship=5] V [marital-status=2^education=11.0-15.0^age=0.41-0.49^workclass=2^occupation=0.0-2.0^hours-per-week=0.38-0.4] V [marital-status=2^education=11.0-15.0^education-num=0.8-1.0] V [marital-status=2^occupation=2.0-3.0^education-num=0.6-0.8^workclass=2^age=0.29-0.34] V [marital-status=2^education-num=0.6-0.8^occupation=9.0-11.0] V [marital-status=2^occupation=2.0-3.0^workclass=2^age=0.49-0.63^education=11.0-15.0] V [marital-status=2^education-

```
num=0.53-0.6^occupation=2.0-3.0^age=0.41-0.49^workclass=3] V [marital-
status=2^education-num=0.53-0.6^capital-loss=0.0-0.98^age=0.34-0.41^hours-per-
week=0.38-0.4] V [marital-status=2^education-num=0.53-0.6^capital-
gain=0.0-1.0^race=2] V [marital-status=2^occupation=2.0-3.0^workclass=2] V
[marital-status=2^capital-gain=0.0-1.0^occupation=9.0-11.0]>
```

```
[21]: print('Accuracy:',ripper_clf.score(X_test, y_test))
print('Pessimistic error: {}'.format(np.divide(np.sum([cmat[0,1],cmat[1,0]]),np.
→sum(cmat))))
```

Accuracy: 0.8872803624709913
Pessimistic error: 0.20035363023538513

```
[22]: #performance of ripper
from sklearn.metrics import precision_score, recall_score
precision = ripper_clf.score(X_test, y_test, precision_score)
recall = ripper_clf.score(X_test, y_test, recall_score)
print(f'precision ripper: {precision} recall ripper: {recall}')
```

precision ripper: 0.0 recall ripper: 0.0

3 Clustering data1 - titanic

```
[23]: #importing data

df_train= pd.read_csv("dataClustering/train.csv")
df_test= pd.read_csv("dataClustering/test.csv")

#preprocessing
df_train = df_train.drop('Name', axis=1,)
df_train = df_train.drop('Ticket', axis=1,)
df_train = df_train.drop('Fare', axis=1,)
df_train = df_train.drop('Cabin', axis=1,)

df_train['Family'] = df_train['SibSp'] + df_train['Parch'] + 1

df_train = df_train.drop('SibSp', axis=1,)
df_train = df_train.drop('Parch', axis=1,)

df_train["Age"] = df_train["Age"].fillna(df_train["Age"].median())
df_train["Embarked"] = df_train["Embarked"].fillna("S")
df_train = df_train.drop('Age', axis=1,)
df_train.head()
```

```
[23]:
```

	PassengerId	Survived	Pclass	Sex	Embarked	Family
0	1	0	3	male	S	2
1	2	1	1	female	C	2
2	3	1	3	female	S	1
3	4	1	1	female	S	2
4	5	0	3	male	S	1

```
[24]: df1 = df_train.filter(['Pclass', 'Sex', 'Embarked', 'Family', 'Adult'], axis=1)

X = df1

df2 = df_train['Survived']

y = df2
```

```
[25]: ### Dropping the Embarked and Family column

X = X.drop('Embarked', axis=1,)
X = X.drop('Family', axis=1,)
```

```
[26]: #testing and splitting data
features_train, features_test, labels_train, labels_test = \
    train_test_split(X, y, test_size=0.3, random_state=42)
```

```
[27]: from sklearn.preprocessing import LabelEncoder

lb_make = LabelEncoder()
features_train['Sex'] = lb_make.fit_transform(features_train['Sex'])
```

```
[28]: from sklearn.preprocessing import LabelEncoder

lb_make = LabelEncoder()
features_test['Sex'] = lb_make.fit_transform(features_test
                                             ['Sex'])
```

```
[29]: # applying 5NN , id3, Naive bayes
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.model_selection import StratifiedKFold, KFold
from statistics import mean, stdev
# visualization
import seaborn as sns
```

```

# Spot-Check Algorithms
models = []

models.append(( 'KNN' , KNeighborsClassifier(n_neighbors = 5, metric = 'minkowski', p = 2)))
models.append(( 'ID3' , DecisionTreeClassifier(criterion = 'entropy', random_state = 0)))
models.append(( 'NB' , GaussianNB()))

# Test options and evaluation metric
num_folds = 5
num_instances = len(features_train)
seed = 7
scoring = 'accuracy'

# Test options and evaluation metric
num_folds = 5
num_instances = len(features_train)
seed = 7
scoring = 'accuracy'
results = []
names = []
print("For Titanic dataset ")
print('model mean standard deviation')
for name, model in models:

    cv_results = StratifiedKFold(n_splits=5, random_state=1)
    scores = cross_val_score(model, X_train, y_train, scoring='accuracy', cv=cv_results, n_jobs=-1)
    model.fit(features_train, labels_train)

    pred = model.predict(features_test)

    results.append(model.score(features_test, labels_test))
    names.append(name)
    msg = "%s: %f (%f)" % (name, scores.mean(), scores.std())
    print('      Mean      Standard Deviation')
    print(msg)
    print('confusion matrix')
    print(confusion_matrix(labels_test, pred))
    # Print out confusion matrix
    cmat = confusion_matrix(labels_test, pred)

    print('TP - True Negative {}'.format(cmat[0,0]))
    print('FP - False Positive {}'.format(cmat[0,1]))

```

```

print('FN - False Negative {}'.format(cmat[1,0]))
print('TP - True Positive {}'.format(cmat[1,1]))
    #accuracy
Accuracy = (np.divide(np.sum([cmat[0,0],cmat[1,1]]),np.sum(cmat)))
print('Accuracy Rate:',Accuracy)
    #interval = z * sqrt( (accuracy * (1 - accuracy)) / n)
    #confidence interval with 80%
confidence_interval = 1.282 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 80% Confidence Interval): ", confidence_interval)
confidence_interval2 = 1.645 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 90% Confidence Interval): ", confidence_interval2)
    #pessimistic error
print('Pessimistic error: {}'.format(np.divide(np.
→sum([cmat[0,1],cmat[1,0]]),np.sum(cmat))))
print('\n')
print()

print('→ 5-Fold cross-validation accuray score for the training data for_
→above classifiers')
print('test results',results)

```

For Titanic dataset

model mean standard deviation

Mean Standard Deviation

KNN: 0.818216 (0.004867)

confusion matrix

[[152 5]

[56 55]]

TP - True Negative 152

FP - False Positive 5

FN - False Negative 56

TP - True Positive 55

Accuracy Rate: 0.7723880597014925

True error (with 80% Confidence Interval): 0.02153571514161661

True error (with 90% Confidence Interval): 0.02763358144146593

Pessimistic error: 0.22761194029850745

Mean Standard Deviation

ID3: 0.810969 (0.009602)

confusion matrix

[[152 5]

[56 55]]

TP - True Negative 152

FP - False Positive 5

FN - False Negative 56

TP - True Positive 55
 Accuracy Rate: 0.7723880597014925
 True error (with 80% Confidence Interval): 0.02153571514161661
 True error (with 90% Confidence Interval): 0.02763358144146593
 Pessimistic error: 0.22761194029850745

Mean Standard Deviation
 NB: 0.804481 (0.006848)
 confusion matrix
 [[134 23]
 [33 78]]
 TP - True Negative 134
 FP - False Positive 23
 FN - False Negative 33
 TP - True Positive 78
 Accuracy Rate: 0.7910447761194029
 True error (with 80% Confidence Interval): 0.020881954819997876
 True error (with 90% Confidence Interval): 0.026794708017859988
 Pessimistic error: 0.208955223880597

-> 5-Fold cross-validation accuray score for the training data for above classifiers
 test results [0.7723880597014925, 0.7723880597014925, 0.7910447761194029]

```
[30]: # applying ripper algorithm
import wittgenstein as lw
ripper_clf = lw.RIPPER() # Or irep_clf = lw.IREP() to build a model using IREP
ripper_clf.fit( features_train,labels_train) # Or pass X and y data to .fit
ripper_clf
```

```
[30]: <RIPPER(dl_allowance=64, k=2, max_rules=None, prune_size=0.33,
random_state=None, verbosity=0, max_rule_conds=None, max_total_conds=None,
n_discretize_bins=10)>
```

```
[31]: # ruleset of ripper algo
ripper_clf.ruleset_
```

```
[31]: <Ruleset [Sex=0^Pclass=1] V [Sex=0^Pclass=2]>
```

```
[32]: # ripper test performance

print("Accuracy:",ripper_clf.score(features_train,labels_train))
```

```
print('Pessimistic error: {}'.format(np.divide(np.sum([cmat[0,1],cmat[1,0]]),np.
↪sum(cmat))))
```

Accuracy: 0.6324237560192616

Pessimistic error: 0.208955223880597

```
[33]: #performance of ripper
from sklearn.metrics import precision_score, recall_score
precision = ripper_clf.score(features_train,labels_train, precision_score)
recall = ripper_clf.score(features_train,labels_train, recall_score)
print(f'precision ripper: {precision}, recall ripper: {recall}')
```

precision ripper: 1.0, recall ripper: 0.008658008658008658

```
[34]: pred1 = ripper_clf.predict(features_test)

print('Confusion Matrix')
print(confusion_matrix(labels_test,pred1))
print('\n')
print("Classification Report")
print(classification_report(labels_test,pred1))
```

Confusion Matrix

```
[[157  0]
 [109  2]]
```

Classification Report

	precision	recall	f1-score	support
0	0.59	1.00	0.74	157
1	1.00	0.02	0.04	111
accuracy			0.59	268
macro avg	0.80	0.51	0.39	268
weighted avg	0.76	0.59	0.45	268

```
[35]: # Print out confusion matrix
cmat = confusion_matrix(labels_test, pred1)
#print(cmat)
print('TP - True Negative {}'.format(cmat[0,0]))
print('FP - False Positive {}'.format(cmat[0,1]))
print('FN - False Negative {}'.format(cmat[1,0]))
print('TP - True Positive {}'.format(cmat[1,1]))
Accuracy = (np.divide(np.sum([cmat[0,0],cmat[1,1]]),np.sum(cmat)))
print('Accuracy Rate:',Accuracy)
```

```

#interval = z * sqrt( (accuracy * (1 - accuracy)) / n)
#confidence interval with 80%
confidence_interval = 1.282 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 80% Confidence Interval): ", confidence_interval)
confidence_interval2 = 1.645 * sqrt( Accuracy * (1 - Accuracy)/num_instances)
print("True error (with 90% Confidence Interval): ", confidence_interval2)
print('Pessimistic Error: {}'.format(np.divide(np.sum([cmat[0,1],cmat[1,0]]),np.
→sum(cmat))))
print('\n')
# Pessimistic error
print('Pessimistic Error: {}'.format(np.divide(np.sum([cmat[0,1],cmat[1,0]]),np.
→sum(cmat))))

```

TP - True Negative 157

FP - False Positive 0

FN - False Negative 109

TP - True Positive 2

Accuracy Rate: 0.5932835820895522

True error (with 80% Confidence Interval): 0.02523021892538764

True error (with 90% Confidence Interval): 0.03237418887071971

Pessimistic Error: 0.40671641791044777

Pessimistic Error: 0.40671641791044777