

Final_file

September 6, 2020

1 Preprocessing of Discrete Data

```
[108]: #importing libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
[109]: # reading data
data = pd.read_csv('dataDiscrete/data.csv')
```

Data Description: This dataset contains characteristics derived from digitized imaging of fine needle aspirates of a breast tumor cell mass. The goal of this analysis is to train a machine learning algorithms to accurately distinguish between a benign and malignant tumor to aid in clinical diagnosis.

Ten features were computed for each cell nucleus:

- a) radius - ratio , real value
- b) texture - ratio ,real
- c) perimeter - ratio ,real
- d) area -ratio , real value
- e) smoothness -ratio , real value
- f) compactness -ratio , real value
- g) concavity -ratio , real value
- h) concave points -ratio , real value
- i) symmetry-ratio , real value
- j) fractal dimension-ratio , real value

<https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Diagnostic%29>

```
[110]: #shape of data
print(data.shape)
# show starting 5 rows
data.head()
```

(569, 33)

```
[110]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

	smoothness_mean	compactness_mean	concavity_mean	concave	points_mean	\
0	0.11840	0.27760	0.3001		0.14710	
1	0.08474	0.07864	0.0869		0.07017	
2	0.10960	0.15990	0.1974		0.12790	
3	0.14250	0.28390	0.2414		0.10520	
4	0.10030	0.13280	0.1980		0.10430	

...	texture_worst	perimeter_worst	area_worst	smoothness_worst	\
0	...	17.33	184.60	2019.0	0.1622
1	...	23.41	158.80	1956.0	0.1238
2	...	25.53	152.50	1709.0	0.1444
3	...	26.50	98.87	567.7	0.2098
4	...	16.67	152.20	1575.0	0.1374

	compactness_worst	concavity_worst	concave	points_worst	symmetry_worst	\
0	0.6656	0.7119		0.2654	0.4601	
1	0.1866	0.2416		0.1860	0.2750	
2	0.4245	0.4504		0.2430	0.3613	
3	0.8663	0.6869		0.2575	0.6638	
4	0.2050	0.4000		0.1625	0.2364	

	fractal_dimension_worst	Unnamed: 32
0	0.11890	NaN
1	0.08902	NaN
2	0.08758	NaN
3	0.17300	NaN
4	0.07678	NaN

[5 rows x 33 columns]

```
[111]: # checking for null data
data.isna().sum()
```

```
[111]: id                                0
      diagnosis                          0
      radius_mean                        0
      texture_mean                       0
      perimeter_mean                     0
      area_mean                          0
      smoothness_mean                    0
      compactness_mean                   0
      concavity_mean                     0
      concave points_mean                 0
      symmetry_mean                       0
      fractal_dimension_mean              0
      radius_se                           0
      texture_se                          0
      perimeter_se                        0
      area_se                             0
      smoothness_se                       0
      compactness_se                      0
      concavity_se                        0
      concave points_se                   0
      symmetry_se                         0
      fractal_dimension_se                0
      radius_worst                       0
      texture_worst                       0
      perimeter_worst                     0
      area_worst                          0
      smoothness_worst                    0
      compactness_worst                   0
      concavity_worst                     0
      concave points_worst                0
      symmetry_worst                      0
      fractal_dimension_worst              0
      Unnamed: 32                         569
      dtype: int64
```

```
[112]: # it shows that all the data are unique that is.
      data['id'].nunique()

      # here data in last column is empty and id is unique, so removing this does not
      → affect data

      data.drop(data.columns[[-1, 0]], axis=1, inplace=True)

      data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
```

```

Data columns (total 31 columns):
diagnosis           569 non-null object
radius_mean        569 non-null float64
texture_mean       569 non-null float64
perimeter_mean     569 non-null float64
area_mean          569 non-null float64
smoothness_mean    569 non-null float64
compactness_mean   569 non-null float64
concavity_mean     569 non-null float64
concave points_mean 569 non-null float64
symmetry_mean      569 non-null float64
fractal_dimension_mean 569 non-null float64
radius_se          569 non-null float64
texture_se         569 non-null float64
perimeter_se       569 non-null float64
area_se           569 non-null float64
smoothness_se      569 non-null float64
compactness_se     569 non-null float64
concavity_se       569 non-null float64
concave points_se  569 non-null float64
symmetry_se        569 non-null float64
fractal_dimension_se 569 non-null float64
radius_worst       569 non-null float64
texture_worst      569 non-null float64
perimeter_worst    569 non-null float64
area_worst         569 non-null float64
smoothness_worst   569 non-null float64
compactness_worst  569 non-null float64
concavity_worst    569 non-null float64
concave points_worst 569 non-null float64
symmetry_worst     569 non-null float64
fractal_dimension_worst 569 non-null float64
dtypes: float64(30), object(1)
memory usage: 137.9+ KB

```

```

[113]: # first 5 entries of dataframe
data.head()

```

```

[113]:  diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0         M        17.99        10.38        122.80        1001.0
1         M        20.57        17.77        132.90        1326.0
2         M        19.69        21.25        130.00        1203.0
3         M        11.42        20.38         77.58         386.1
4         M        20.29        14.34        135.10        1297.0

    smoothness_mean  compactness_mean  concavity_mean  concave points_mean  \
0         0.11840         0.27760         0.3001         0.14710

```

1	0.08474	0.07864	0.0869	0.07017
2	0.10960	0.15990	0.1974	0.12790
3	0.14250	0.28390	0.2414	0.10520
4	0.10030	0.13280	0.1980	0.10430

	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	\
0	0.2419	...	25.38	17.33	184.60	
1	0.1812	...	24.99	23.41	158.80	
2	0.2069	...	23.57	25.53	152.50	
3	0.2597	...	14.91	26.50	98.87	
4	0.1809	...	22.54	16.67	152.20	

	area_worst	smoothness_worst	compactness_worst	concavity_worst	\
0	2019.0	0.1622	0.6656	0.7119	
1	1956.0	0.1238	0.1866	0.2416	
2	1709.0	0.1444	0.4245	0.4504	
3	567.7	0.2098	0.8663	0.6869	
4	1575.0	0.1374	0.2050	0.4000	

	concave points_worst	symmetry_worst	fractal_dimension_worst
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678

[5 rows x 31 columns]

[114]: *#finding out numerical and categorical*

```
numerical_features=[feature for feature in data.columns if data[feature].dtype!
    ↳='O']
categorical_features=[feature for feature in data.columns if data[feature].
    ↳dtype=='O' and feature!='wage_class']

print('categorical features: '\n',categorical_features)
print('\n')
print('numerical features: '\n',numerical_features)
```

categorical features:
['diagnosis']

numerical features:
['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean',
'smoothness_mean', 'compactness_mean', 'concavity_mean', 'concave points_mean',
'symmetry_mean', 'fractal_dimension_mean', 'radius_se', 'texture_se',

```
'perimeter_se', 'area_se', 'smoothness_se', 'compactness_se', 'concavity_se',
'concave points_se', 'symmetry_se', 'fractal_dimension_se', 'radius_worst',
'texture_worst', 'perimeter_worst', 'area_worst', 'smoothness_worst',
'compactness_worst', 'concavity_worst', 'concave points_worst',
'symmetry_worst', 'fractal_dimension_worst']
```

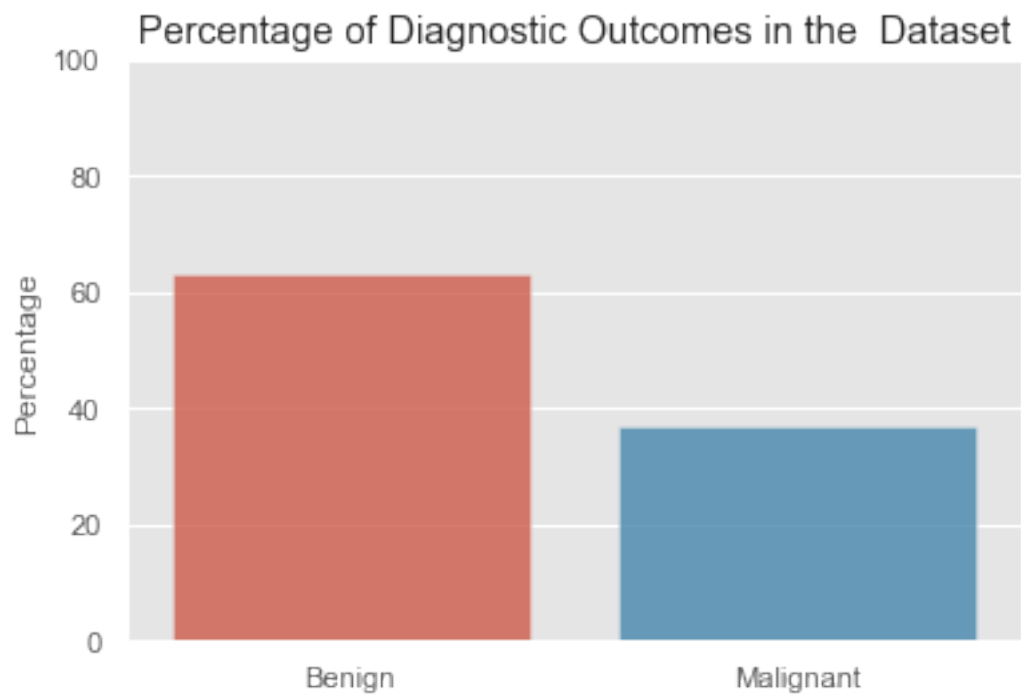

```
[115]: data_count=data.diagnosis.value_counts(normalize = True)
data_count = pd.Series(data_count)
data_count = pd.DataFrame(data_count)
data_count.index = ['Benign', 'Malignant']

data_count['Percent'] = 100*data_count['diagnosis']/sum(data_count['diagnosis'])
data_count['Percent'] = data_count['Percent'].round().astype('int')
data_count
```

```
[115]:
```

	diagnosis	Percent
Benign	0.627417	63
Malignant	0.372583	37

```
[116]: # Visualize frequency and percentage of Diagnostic Outcomes in the Dataset
sns.barplot(x = ['Benign', 'Malignant'], y = 'diagnosis', data = data_count,
            alpha = .8)
plt.title('Frequency of Diagnostic Outcomes in the Dataset')
plt.ylabel('Frequency')
plt.show()
sns.barplot(x = ['Benign', 'Malignant'], y = 'Percent', data = data_count,
            alpha = .8)
plt.title('Percentage of Diagnostic Outcomes in the Dataset')
plt.ylabel('Percentage')
plt.ylim(0,100)
plt.show()
```



```
[117]: # changing categorical data to numerical
data['diagnosis'] = data['diagnosis'].map({'M':1, 'B':0})
```

```
[118]: # checking the different values contained in the diagnosis column
#Benign : 0
#Malign : 1

data['diagnosis'].value_counts()
```

```
[118]: 0    357
      1    212
      Name: diagnosis, dtype: int64
```

```
[119]: # 'duplicated()' function in pandas return the duplicate row as True and other
      ↪ as False
#for counting the duplicate elements we sum all the rows
sum(data.duplicated())
```


```
[119]: 0
```

```
[120]: # checking data shape again.
data.shape
```

```
[120]: (569, 31)
```

```
[121]: # Finding correlation among features using sns' heatmap
plt.figure(figsize=(25,25))
sns.heatmap(data.corr(),annot=True, cmap="YlGnBu")
```

```
[121]: <matplotlib.axes._subplots.AxesSubplot at 0x299366f7dc8>
```

Observation: -> The radius, parameter and area are highly correlated as expected from their relation so from these we will use anyone of them

-> Compactness_mean, concavity_mean and concavepoint_mean are highly correlated so we will use compactness_mean from here

-> So selected Parameter for use is perimeter_mean, texture_mean, compactness_mean, symmetry_mean

Dimensionality Reduction From 30 components which are the most important ones (interestingness)? Can we reduce our data dimension? Approach:

1. scaling

2. PCA

Essentially the same process for each of the above two steps:

- i) import
- ii) instantiate
- iii) fit
- iv) transform

```
[122]: ### Applying Dimensionality Reduction

from sklearn.preprocessing import StandardScaler # Import
scaler = StandardScaler() # Instantiate
scaler.fit(data) # Fit
scaled_data = scaler.transform(data) # Transform

# Applying PCA
from sklearn.decomposition import PCA # Import

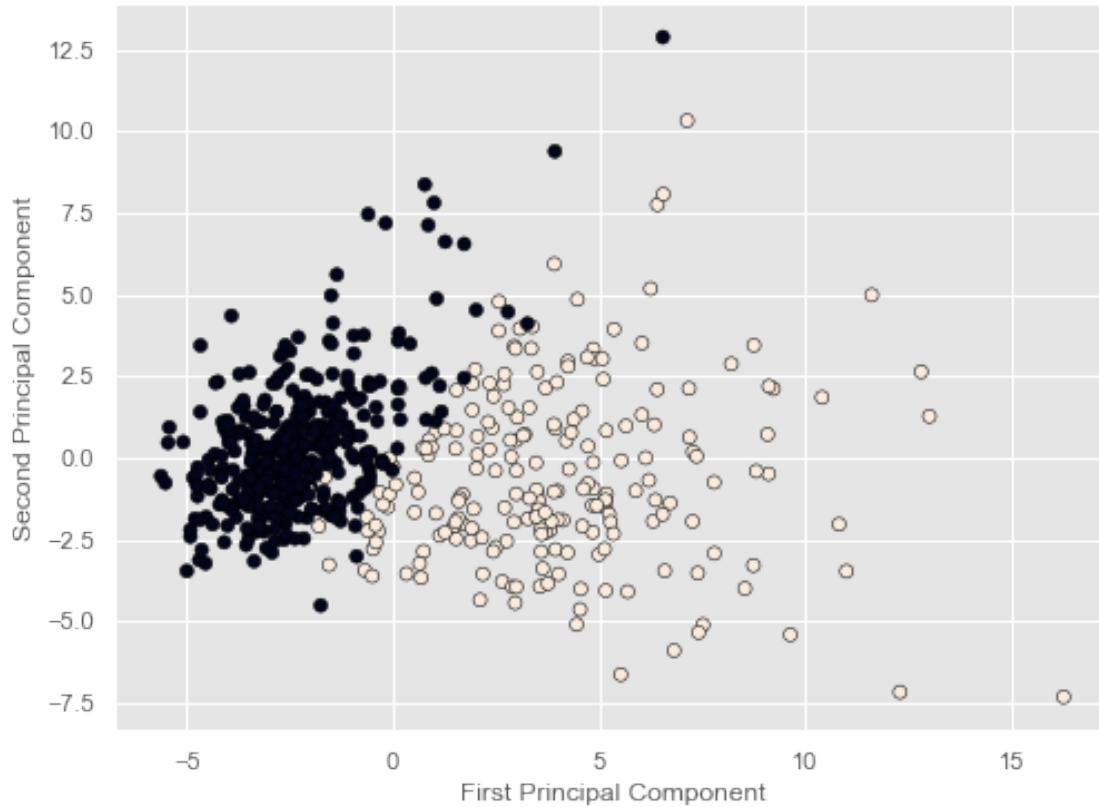
pca = PCA(n_components=2) # Instantiate
pca.fit(scaled_data) # Fit
X_pca = pca.transform(scaled_data) # Transform
print(" Data dimensions before reduction:",scaled_data.shape)
print(" Data dimensions after reduction:",X_pca.shape)

# visualization of PCA
plt.figure(figsize=(8,6))
plt.scatter(X_pca[:,0],X_pca[:,1], c=data['diagnosis'],edgecolor='black')
#plt.scatter(X_pca[:,0],X_pca[:,1], edgecolor='black')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
```

```
Data dimensions before reduction: (569, 31)
```

```
Data dimensions after reduction: (569, 2)
```

```
[122]: Text(0, 0.5, 'Second Principal Component')
```



```
[123]: print(pca.components_)
```

```
[[ 0.21691695  0.21639914  0.10359936  0.2245478   0.21796464  0.13764549
  0.23150388  0.25122179  0.2552664   0.1330126   0.05797189  0.20090409
  0.01547414  0.20563036  0.19798194  0.01123624  0.1619218   0.14578435
  0.1762679   0.03832541  0.09508414  0.22558983  0.10501867  0.23364163
  0.22196962  0.12530106  0.20447639  0.22307483  0.24628844  0.12066328
  0.12672101]
[-0.07760994 -0.22654473 -0.05826196 -0.20762951 -0.22322434  0.18876038
  0.15847307  0.06786768 -0.02722028  0.19321967  0.36761956 -0.09715581
  0.09118826 -0.08069084 -0.14355563  0.20531469  0.23886735  0.203226
  0.13704859  0.1860032   0.28433173 -0.21289639 -0.04513536 -0.19259202
 -0.21188718  0.17247591  0.14766283  0.10308807 -0.00243309  0.14206245
  0.27644912]]
```

```
[124]: print(pca.explained_variance_)
```

```
[13.94227406  5.73643378]
```

```
[125]: #checking normalization with mean and standard deviation
print('Mean after applying PCA', np.mean(X_pca))
```

```
print('Standard Deviation after applying PCA',np.std(X_pca))

from scipy.stats import kurtosis
print('Kurtosis after applying PCA',kurtosis(X_pca))
```

Mean after applying PCA -2.4975140097015822e-17
Standard Deviation after applying PCA 3.1340168409862255
Kurtosis after applying PCA [0.6508818 2.857848]

Data is in normal distribution after applying PCA here. No need to further process.

```
[126]: # converting the normalized features into a tabular format with the help of
↳ DataFrame.
pca_Df = pd.DataFrame(data = X_pca
                      , columns = ['principal component 1', 'principal component 2'])
pca_Df.head()
```

```
[126]:    principal component 1  principal component 2
0                9.225770                2.116196
1                2.655802               -3.784776
2                5.892492               -1.005579
3                7.135401               10.318716
4                4.129423               -1.905579
```

```
[127]: pca_Df.corr()
```

```
[127]:    principal component 1  principal component 2
principal component 1    1.000000e+00    -8.882348e-17
principal component 2   -8.882348e-17    1.000000e+00
```

High complexity associated with dataframe having a big number of dimensions/features, which frequently make the target function quite complex and may lead to model overfitting as long as often the dataset lies on the lower dimensionality manifold.

As PCA convert more number of features columns to 2 dimensional data, we have output for 2 columns as principal component 1 and principal component 2.

So that we have reduced set of correlation between attributes,

```
[128]: # Shuffling the Data Set
from sklearn.utils import shuffle
X_pca = shuffle(X_pca)

# Splitting the data set into train and test set
from sklearn.model_selection import train_test_split

features = data.drop(columns = ['diagnosis'])
target = data['diagnosis']
```

```
X_train1, X_test1, y_train1, y_test1 = train_test_split(features, target,
↳test_size = 0.3,random_state = 0)

print ("Train data set size : ", X_train1.shape)
print ("Test data set size : ", X_test1.shape)
```


```
Train data set size : (398, 30)
Test data set size : (171, 30)
```

Feature importance A benefit of using gradient boosting is that after the boosted trees are constructed, it is relatively straightforward to retrieve importance scores for each attribute.

```
[129]: # Plotting the feature importances using the Boosted Gradient Descent
from xgboost import XGBClassifier
from xgboost import plot_importance

# Training the model
model = XGBClassifier()
model_importance = model.fit(X_train1, y_train1)

# Plotting the Feature importance bar graph
plt.rcParams['figure.figsize'] = [14,12]
sns.set(style = 'darkgrid')
plot_importance(model_importance);
```



2 Preprocessing of Continuous Data

Adult / Census Income dataset Data can be downloaded from following link
<https://archive.ics.uci.edu/ml/datasets/census+income>

```
[130]: # Import libraries necessary for this project
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import seaborn as sns
sns.set(style="darkgrid")
from time import time

from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import OneHotEncoder
import scipy.stats as stats
from scipy.stats import kurtosistest

# displaying for notebooks
```

```
%matplotlib inline

columns = ['age', 'workclass', 'fnlwgt', 'education', 'education-num',
           ↪ 'marital-status', 'occupation',
           ↪ 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss',
           ↪ 'hours-per-week', 'native-country', 'income']

# Load the Census dataset
census = pd.read_csv('dataConti/adult.data', header=None, names=columns,
                    ↪ skipinitialspace=True)

# Success - Display the first 5 record
display(census.head())
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
0	Never-married	Adm-clerical	Not-in-family	White	Male	
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
2	Divorced	Handlers-cleaners	Not-in-family	White	Male	
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

```
[131]: # checking shape of data set
census.shape
```

```
[131]: (32561, 15)
```

```
[132]: import warnings
warnings.filterwarnings("ignore")
```

Data Information age: continuous.

workclass: categorical (Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov,

Without-pay, Never-worked)

fnlwgt: continuous.

education: Categorical (Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool)

education-num: continuous.

marital-status: Categorical

relationship: Categorical

race: Categorical

sex: Female, Male - categorical

capital-gain: continuous (income from investment sources, apart from wages/salary)

capital-loss: continuous (losses from investment sources, apart from wages/salary)

hours-per-week: continuous.

native-country: categorical.

```
[133]: #finding out numerical and categorical

numerical_features=[feature for feature in census.columns if census[feature].
    ↳dtype!='O']
categorical_features=[feature for feature in census.columns if 
    ↳census[feature].dtype=='O' and feature!='wage_class']

print('categorical features: '\n',categorical_features)
print('\n')
print('numerical features: '\n',numerical_features)
```

categorical features:

['workclass', 'education', 'marital-status', 'occupation', 'relationship',
'race', 'sex', 'native-country', 'income']

numerical features:

['age', 'fnlwgt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-
week']

```
[134]: # Examine if there are missing value
census.isna().sum()
```

```
[134]: age                0
workclass              0
fnlwgt                 0
education              0
education-num          0
```



```

marital-status    0
occupation        0
relationship      0
race              0
sex              0
capital-gain      0
capital-loss      0
hours-per-week    0
native-country    0
income            0
dtype: int64

```

The result above shows there's no null value in dataset. But according to data notes provided, unknown data was converted into '?'. Therefore, next we'll convert '?' to NaNs and drop unwanted columns.

```

[135]: # Drop the fnlwgt column which is useless for later analysis
census = census.drop('fnlwgt', axis=1)

# Read in test data
census_test = pd.read_csv('dataConti/adult.test', header=None, skiprows=1,
↳names=columns, skipinitialspace=True)

# Drop the fnlwgt column which is useless for later analysis
census_test = census_test.drop('fnlwgt', axis=1)

# Remove '.' in income column
census_test['income'] = census_test['income'].apply(lambda x: '>50K' if
↳x=='>50K.' else '<=50K')

# Review several rows and shape of data set
display(census_test.head())
display(census_test.shape)

```

	age	workclass	education	education-num	marital-status	\
0	25	Private	11th	7	Never-married	
1	38	Private	HS-grad	9	Married-civ-spouse	
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	
3	44	Private	Some-college	10	Married-civ-spouse	
4	18	?	Some-college	10	Never-married	

	occupation	relationship	race	sex	capital-gain	capital-loss	\
0	Machine-op-inspct	Own-child	Black	Male	0	0	
1	Farming-fishing	Husband	White	Male	0	0	
2	Protective-serv	Husband	White	Male	0	0	
3	Machine-op-inspct	Husband	Black	Male	7688	0	
4	?	Own-child	White	Female	0	0	

	hours-per-week	native-country	income
0	40	United-States	<=50K
1	50	United-States	<=50K
2	40	United-States	>50K
3	40	United-States	>50K
4	30	United-States	<=50K

(16281, 14)

```
[136]: # Convert '?' to NaNs and remove the entries with NaN value
object_col = census.select_dtypes(include=object).columns.tolist()
for col in object_col:
    census.loc[census[col]=='?', col] = np.nan
    census_test.loc[census_test[col]=='?', col] = np.nan

# Perform an mssing assessment in each column of the dataset.
col_missing_pct = census.isna().sum()/census.shape[0]
col_missing_pct.sort_values(ascending=False)
```

```
[136]: occupation      0.056601
workclass      0.056386
native-country  0.017905
income         0.000000
hours-per-week 0.000000
capital-loss    0.000000
capital-gain    0.000000
sex            0.000000
race           0.000000
relationship    0.000000
marital-status 0.000000
education-num   0.000000
education       0.000000
age            0.000000
dtype: float64
```

The largest missing percentage by column level is 5% in dataset, and most columns are complete enough. Therefore, here I'll remove the NaN values instead of manually imputing.

```
[137]: # Removing data entries with missing value
adult_train = census.dropna(axis=0, how='any')
adult_test = census_test.dropna(axis=0, how='any')

# Show the results of the split
print("After removing the missing value:")
print("Training set has {} samples.".format(adult_train.shape[0]))
print("Testing set has {} samples.".format(adult_test.shape[0]))
```

After removing the missing value:
 Training set has 30162 samples.
 Testing set has 15060 samples.

```
[138]: # Finding correlation among features
census.corr() #before removing missing values
```

```
[138]:
```

	age	education-num	capital-gain	capital-loss	\
age	1.000000	0.036527	0.077674	0.057775	
education-num	0.036527	1.000000	0.122630	0.079923	
capital-gain	0.077674	0.122630	1.000000	-0.031615	
capital-loss	0.057775	0.079923	-0.031615	1.000000	
hours-per-week	0.068756	0.148123	0.078409	0.054256	


```

hours-per-week
age      0.068756
education-num  0.148123
capital-gain  0.078409
capital-loss  0.054256
hours-per-week  1.000000

```

We can see correlation changes after removing missing values.

Correlation after removing missing values changes. This means the change in one variable reflects a change in another variable in a predictable pattern then we say that the variables are correlated.

Here it is observable that all are positively correlated except capital-loss and capital-gain as it is negatively correlated.

```
[139]: adult_train.corr() # after removing missing values
```

```
[139]:
```

	age	education-num	capital-gain	capital-loss	\
age	1.000000	0.043526	0.080154	0.060165	
education-num	0.043526	1.000000	0.124416	0.079646	
capital-gain	0.080154	0.124416	1.000000	-0.032229	
capital-loss	0.060165	0.079646	-0.032229	1.000000	
hours-per-week	0.101599	0.152522	0.080432	0.052417	


```

hours-per-week
age      0.101599
education-num  0.152522
capital-gain  0.080432
capital-loss  0.052417
hours-per-week  1.000000

```

```
[140]: # Check the skewness of numerical variables in data set
num_col = adult_train.dtypes[adult_train.dtypes != 'object'].index

# Create figure
```

```

fig = plt.figure(figsize = (10,13));

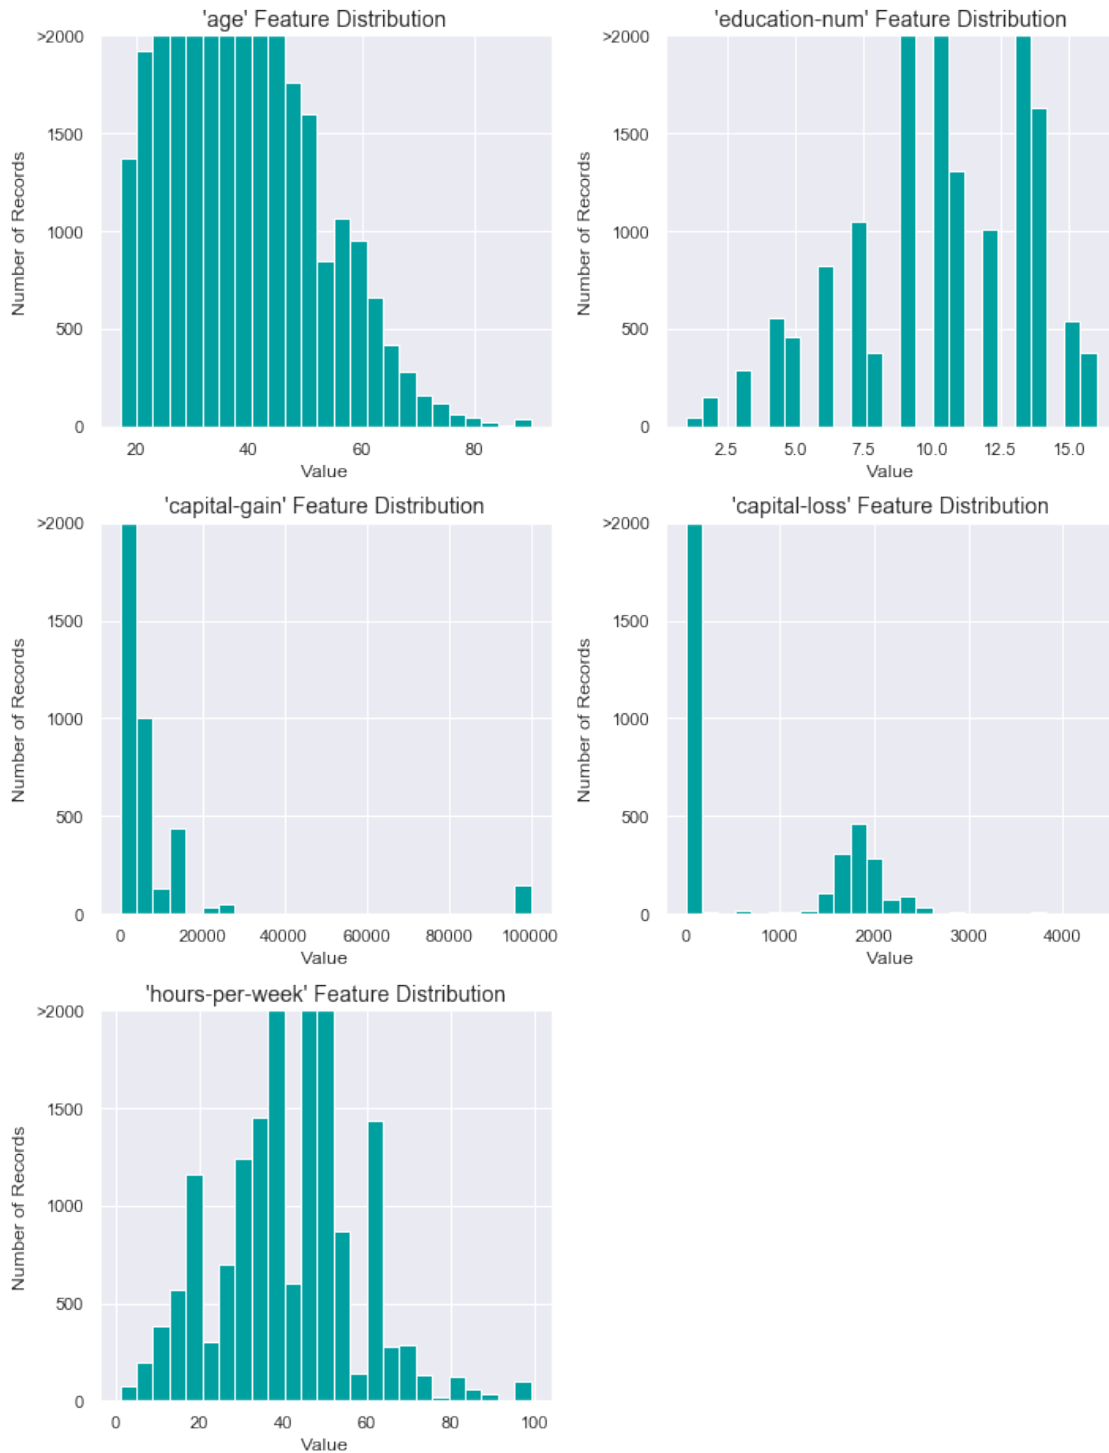
# Skewed feature plotting
for i, feature in enumerate(adult_train[num_col]):
    ax = fig.add_subplot(3, 2, i+1)
    ax.hist(adult_train[feature], bins = 25, color = '#00A0A0')
    ax.set_title("%s' Feature Distribution"%(feature), fontsize = 14)
    ax.set_xlabel("Value")
    ax.set_ylabel("Number of Records")
    ax.set_ylim((0, 2000))
    ax.set_yticks([0, 500, 1000, 1500, 2000])
    ax.set_yticklabels([0, 500, 1000, 1500, ">2000"])

# Plot aesthetics
fig.suptitle("Skewed Distributions of Continuous Census Data Features",
    ↳ fontsize = 16, y = 1.03)

fig.tight_layout()
fig.show()

```

Skewed Distributions of Continuous Census Data Features



As shown in the graph, there seems skewness in 'capital-gain' and 'capital-loss' features. Use quantitative result to confirm if I need to transform skewness in these two variables.

```
[141]: # Calculate skew and sort
skew_feats = adult_train[num_col].skew().sort_values(ascending=False)
skewness = pd.DataFrame({'Skew': skew_feats})
skewness
```

```
[141]:
```

	Skew
capital-gain	11.902682
capital-loss	4.526380
age	0.530228
hours-per-week	0.330869
education-num	-0.305379

Normalizing highly skewed data

```
[142]: # Split the data into features and target label
income_raw = adult_train['income']
feature_raw = adult_train.drop('income', axis=1)

income_raw_test = adult_test['income']
feature_raw_test = adult_test.drop('income', axis=1)

# Log transform the skewed feature highly-skewed feature 'capital-gain' and
# 'capital-loss'.
skewed = ['capital-gain', 'capital-loss']
census_log = pd.DataFrame(data=feature_raw)
census_log[skewed] = feature_raw[skewed].apply(lambda x: np.log(x + 1))

census_log_test = pd.DataFrame(data=feature_raw_test)
census_log_test[skewed] = feature_raw_test[skewed].apply(lambda x: np.log(x +
# 1))

# Initialize a scaler, then apply it to the features
scaler = MinMaxScaler() # default=(0, 1)

features_log_minmax_transform = pd.DataFrame(data = census_log)
features_log_minmax_transform[num_col] = scaler.
# fit_transform(census_log[num_col])

# Transform the test data set
features_log_minmax_transform_test = pd.DataFrame(data = census_log_test)
features_log_minmax_transform_test[num_col] = scaler.
# transform(census_log_test[num_col])

# Show an example of a record with scaling applied
display(features_log_minmax_transform.head())
display(features_log_minmax_transform_test.head())
```

	age	workclass	education	education-num	marital-status \
0	0.301370	State-gov	Bachelors	0.800000	Never-married
1	0.452055	Self-emp-not-inc	Bachelors	0.800000	Married-civ-spouse
2	0.287671	Private	HS-grad	0.533333	Divorced
3	0.493151	Private	11th	0.400000	Married-civ-spouse
4	0.150685	Private	Bachelors	0.800000	Married-civ-spouse

	occupation	relationship	race	sex	capital-gain \
0	Adm-clerical	Not-in-family	White	Male	0.667492
1	Exec-managerial	Husband	White	Male	0.000000
2	Handlers-cleaners	Not-in-family	White	Male	0.000000
3	Handlers-cleaners	Husband	Black	Male	0.000000
4	Prof-specialty	Wife	Black	Female	0.000000

	capital-loss	hours-per-week	native-country
0	0.0	0.397959	United-States
1	0.0	0.122449	United-States
2	0.0	0.397959	United-States
3	0.0	0.397959	United-States
4	0.0	0.397959	Cuba

	age	workclass	education	education-num	marital-status \
0	0.109589	Private	11th	0.400000	Never-married
1	0.287671	Private	HS-grad	0.533333	Married-civ-spouse
2	0.150685	Local-gov	Assoc-acdm	0.733333	Married-civ-spouse
3	0.369863	Private	Some-college	0.600000	Married-civ-spouse
5	0.232877	Private	10th	0.333333	Never-married

	occupation	relationship	race	sex	capital-gain	capital-loss \
0	Machine-op-inspct	Own-child	Black	Male	0.000000	0.0
1	Farming-fishing	Husband	White	Male	0.000000	0.0
2	Protective-serv	Husband	White	Male	0.000000	0.0
3	Machine-op-inspct	Husband	Black	Male	0.777174	0.0
5	Other-service	Not-in-family	White	Male	0.000000	0.0

	hours-per-week	native-country
0	0.397959	United-States
1	0.500000	United-States
2	0.397959	United-States
3	0.397959	United-States
5	0.295918	United-States

```
[143]: census_log.corr()
```

```
[143]:
```

	capital-gain	capital-loss	hours-per-week
capital-gain	1.000000	-0.067040	0.086243

capital-loss	-0.067040	1.000000	0.049468
hours-per-week	0.086243	0.049468	1.000000

```
[144]: census.corr()
```

```
[144]:
```

	age	education-num	capital-gain	capital-loss \
age	1.000000	0.036527	0.077674	0.057775
education-num	0.036527	1.000000	0.122630	0.079923
capital-gain	0.077674	0.122630	1.000000	-0.031615
capital-loss	0.057775	0.079923	-0.031615	1.000000
hours-per-week	0.068756	0.148123	0.078409	0.054256

	hours-per-week
age	0.068756
education-num	0.148123
capital-gain	0.078409
capital-loss	0.054256
hours-per-week	1.000000

Applying log transformation changes correlation. In this case the two coefficients may lead to different statistical inference. For example, a correlation coefficient of 0.2 is considered to be negligible correlation while a correlation coefficient of 0.3 is considered as low positive correlation.

```
[145]: X = feature_raw
       y =income_raw
```

```
[146]: X.head()
```

```
[146]:
```

	age	workclass	education	education-num	marital-status \
0	0.301370	State-gov	Bachelors	0.800000	Never-married
1	0.452055	Self-emp-not-inc	Bachelors	0.800000	Married-civ-spouse
2	0.287671	Private	HS-grad	0.533333	Divorced
3	0.493151	Private	11th	0.400000	Married-civ-spouse
4	0.150685	Private	Bachelors	0.800000	Married-civ-spouse

	occupation	relationship	race	sex	capital-gain \
0	Adm-clerical	Not-in-family	White	Male	0.667492
1	Exec-managerial	Husband	White	Male	0.000000
2	Handlers-cleaners	Not-in-family	White	Male	0.000000
3	Handlers-cleaners	Husband	Black	Male	0.000000
4	Prof-specialty	Wife	Black	Female	0.000000

	capital-loss	hours-per-week	native-country
0	0.0	0.397959	United-States
1	0.0	0.122449	United-States
2	0.0	0.397959	United-States

3	0.0	0.397959	United-States
4	0.0	0.397959	Cuba

```
[147]: # splitting data to train and test for further preprocessing

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
    random_state = 0)
```

```
[148]: from sklearn import preprocessing

categorical = ['workclass', 'education', 'marital-status', 'occupation',
    'relationship', 'race', 'sex', 'native-country']
for feature in categorical:
    le = preprocessing.LabelEncoder()
    X_train[feature] = le.fit_transform(X_train[feature])
    X_test[feature] = le.transform(X_test[feature])
```

```
[149]: # changing data format to pandas.
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train = pd.DataFrame(scaler.fit_transform(X_train), columns = X.columns)

X_test = pd.DataFrame(scaler.transform(X_test), columns = X.columns)
X_train.head()
```

```
[149]:
```

	age	workclass	education	education-num	marital-status	occupation	\
0	0.417684	-0.204511	0.170452	-0.434574	-1.721729	1.251744	
1	-0.494550	-0.204511	1.222721	-0.042490	-0.387728	1.748265	
2	-0.646589	-1.261772	0.170452	-0.434574	-1.721729	1.003483	
3	-1.178726	-0.204511	1.222721	-0.042490	-0.387728	-0.982600	
4	1.177880	-1.261772	1.222721	-0.042490	-0.387728	1.003483	

	relationship	race	sex	capital-gain	capital-loss	\
0	-0.266868	0.382579	-1.438780	-0.301995	4.344011	
1	-0.890070	0.382579	0.695033	-0.301995	-0.223788	
2	0.979538	0.382579	-1.438780	-0.301995	-0.223788	
3	-0.890070	0.382579	0.695033	-0.301995	-0.223788	
4	-0.890070	0.382579	0.695033	-0.301995	-0.223788	

	hours-per-week	native-country
0	-0.078362	0.265946
1	1.605395	0.265946
2	-0.751865	0.265946

3	1.605395	0.265946
4	0.931892	0.265946

```
[150]: # applying dimensionality reduction

from sklearn.decomposition import PCA
# Apply PCA to the data
pca = PCA()
model_adult = pca.fit_transform(X_train)
```

```
[151]: vals = pca.explained_variance_ratio_
vals
```

```
[151]: array([0.16053926, 0.10824834, 0.08848228, 0.08405468, 0.08150518,
          0.07643645, 0.07496046, 0.06805456, 0.06512367, 0.06429608,
          0.05238639, 0.04632909, 0.02958357])
```

```
[152]: # Shuffling the Data Set
from sklearn.utils import shuffle
X_adult_pca = shuffle(model_adult)

# Splitting the data set into train and test set
from sklearn.model_selection import train_test_split


print ("Train data set size : ", X_train.shape)
print ("Test data set size : ", X_test.shape)
```

```
Train data set size : (21113, 13)
Test data set size : (9049, 13)
```

```
[153]: # Plotting the feature importances using the Boosted Gradient Descent
from xgboost import XGBClassifier
from xgboost import plot_importance

# Training the model
model = XGBClassifier()
model_importance_adult = model.fit(X_train, y_train)

# Plotting the Feature importance bar graph
plt.rcParams['figure.figsize'] = [14,12]
sns.set(style = 'darkgrid')
plot_importance(model_importance_adult);
```



3 Transactional Dataset

Data used here is from kaggle can be downloaded from

<https://www.kaggle.com/roshansharma/market-basket-optimization>

Csv file contains information about Customers buying different grocery items at a Mall.

```
[154]: #reading data from csv
data3 = pd.read_csv("dataTransactional/store_data.csv", header=None)

print(data3.shape)
```

(7501, 20)

```
[155]: # checking first 5 entries
data3.head()
```

```
[155]:      0      1      2      3      4  \
0    shrimp almonds avocado vegetables mix green grapes
```

1	burgers	meatballs	eggs	NaN	NaN
2	chutney	NaN	NaN	NaN	NaN
3	turkey	avocado	NaN	NaN	NaN
4	mineral water	milk	energy bar	whole wheat rice	green tea

	5	6	7	8	9 \
0	whole weat flour	yams	cottage cheese	energy drink	tomato juice
1	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN

	10	11	12	13	14	15 \
0	low fat yogurt	green tea	honey	salad	mineral water	salmon
1	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN

	16	17	18	19
0	antioxydant juice	frozen smoothie	spinach	olive oil
1	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

```
[156]: # checking dataframe information
data3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7501 entries, 0 to 7500
Data columns (total 20 columns):
0      7501 non-null object
1      5747 non-null object
2      4389 non-null object
3      3345 non-null object
4      2529 non-null object
5      1864 non-null object
6      1369 non-null object
7      981 non-null object
8      654 non-null object
9      395 non-null object
10     256 non-null object
11     154 non-null object
12     87 non-null object
13     47 non-null object
14     25 non-null object
```

```
15    8 non-null object
16    4 non-null object
17    4 non-null object
18    3 non-null object
19    1 non-null object
dtypes: object(20)
memory usage: 1.1+ MB
```

```
[157]: #missing values
```

```
data3.isna().sum()
```

```
[157]: 0      0
      1    1754
      2    3112
      3    4156
      4    4972
      5    5637
      6    6132
      7    6520
      8    6847
      9    7106
     10    7245
     11    7347
     12    7414
     13    7454
     14    7476
     15    7493
     16    7497
     17    7497
     18    7498
     19    7500
dtype: int64
```

```
[158]: #creating list of itemsets in basket in different transactions
```

```
basket_items = []
for index, row in data3.iterrows():
    cleansed_items = [item for item in row if str(item)!='nan']
    #print(f'basket size: {len(cleansed_items)}, basket:\n{cleansed_items}')
    basket_items.append(cleansed_items)

basket_items[:3]
```

```
[158]: [['shrimp',
        'almonds',
        'avocado',
```

```

'vegetables mix',
'green grapes',
'whole weat flour',
'yams',
'cottage cheese',
'energy drink',
'tomato juice',
'low fat yogurt',
'green tea',
'honey',
'salad',
'mineral water',
'salmon',
'antioxydant juice',
'frozen smoothie',
'spinach',
'olive oil'],
['burgers', 'meatballs', 'eggs'],
['chutney']]

```

[159]: *#Creating transaction DataFrame - this will helps to treate missing values in the dataset.*

```

from mlxtend.preprocessing import TransactionEncoder
tran_encod = TransactionEncoder()
tran_encod_list = tran_encod.fit(basket_items).transform(basket_items)
transaction_df = pd.DataFrame(tran_encod_list, columns=tran_encod.columns_)
transaction_df.head()

```

[159]:

	asparagus	almonds	antioxydant juice	asparagus	avocado	babies food	\
0	False	True	True	False	True	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	True	False	
4	False	False	False	False	False	False	

	bacon	barbecue sauce	black tea	blueberries	...	turkey	vegetables mix	\
0	False	False	False	False	...	False	True	
1	False	False	False	False	...	False	False	
2	False	False	False	False	...	False	False	
3	False	False	False	False	...	True	False	
4	False	False	False	False	...	False	False	

	water spray	white wine	whole weat flour	whole wheat pasta	\
0	False	False	True	False	
1	False	False	False	False	
2	False	False	False	False	
3	False	False	False	False	

4	False	False	False	False
	whole wheat rice	yams	yogurt cake	zucchini
0	False	True	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	True	False	False	False

[5 rows x 120 columns]

Above code have solved the problem of missing values

```
[160]: #checking for missing values
transaction_df.isna().sum()
```

```
[160]: asparagus      0
almonds              0
antioxydant juice    0
asparagus            0
avocado              0
..
whole wheat pasta    0
whole wheat rice     0
yams                 0
yogurt cake          0
zucchini             0
Length: 120, dtype: int64
```

```
[161]: # creating data frame for item frequency
item_count = {}
for col in transaction_df.columns:
    item_count[col] = transaction_df[col].sum()

item_freq_df = pd.DataFrame(data=list(item_count.values()),
    ↪index=list(item_count.keys()), columns=['frequency']).
    ↪sort_values(by='frequency', ascending=False)
item_freq_df.shape, item_freq_df.head(10)
```

```
[161]: ((120, 1),          frequency
mineral water      1788
eggs               1348
spaghetti          1306
french fries       1282
chocolate          1229
green tea          991
milk               972
```

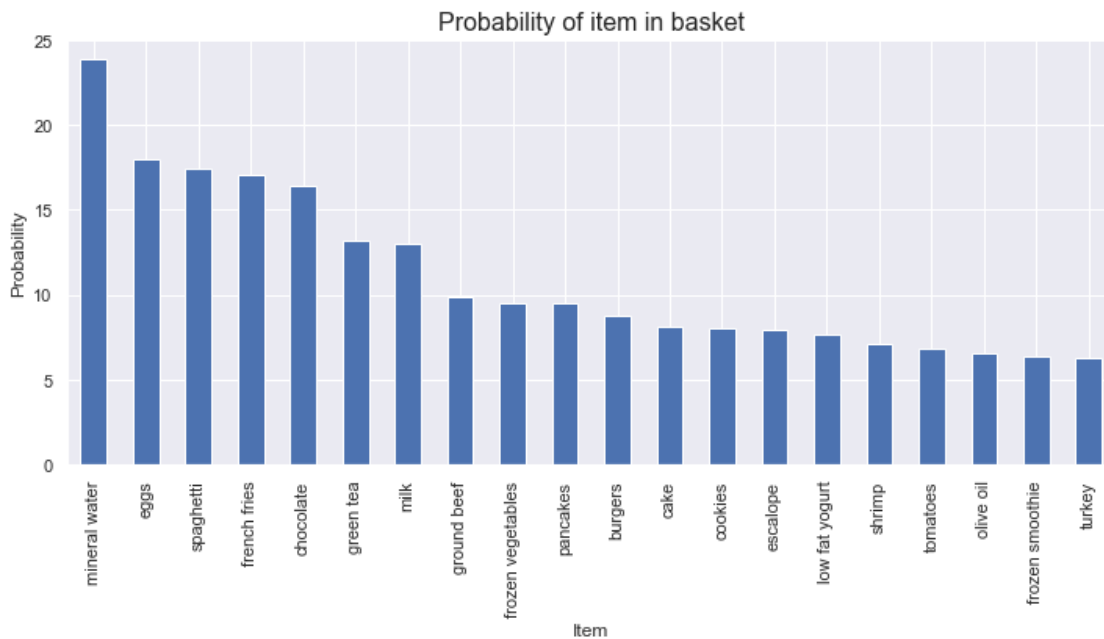
ground beef	737
frozen vegetables	715
pancakes	713)

Frequency shows the sold item in order of most sold to least sold in the store.

```
[162]: #Most purchased items
support =transaction_df.mean()
support.sort_values()

(100*support.sort_values(ascending=False))[:20] .
    ↳plot(kind='bar',grid=True,figsize = (12,5))
plt.title("Probability of item in basket",fontsize = 16)
plt.xlabel('Item')
plt.ylabel('Probability')
```

```
[162]: Text(0, 0.5, 'Probability')
```



4 Clustering 1 - Normal

Dataset is available on kaggle.

Titanic Dataset - Prediction of survived.

<https://www.kaggle.com/c/titanic/data>

Dataset Information: survival binary pclass - (Ticket class 1 = 1st, 2 = 2nd, 3 = 3rd) categorical

sex - categorical,nominal

Age - continous, ratio

sibsp # of siblings / spouses aboard the Titanic - nominal

parch # of parents / children aboard the Titanic - nominal

ticket - nominal

fare - ordinal

cabin - nominal

embarked (Port of Embarkation C = Cherbourg, Q = Queenstown, S = Southampton) - nominal

[163]: *#importing data*

```
df= pd.read_csv("dataClustering/train.csv")
df_test= pd.read_csv("dataClustering/test.csv")
df.head()
```

[163]:

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	


	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

[164]: *#### Visualizing Survival based on gender*

[165]: `df['Died'] = 1 - df['Survived']`

```
df.groupby('Sex').agg('mean')[['Survived', 'Died']].plot(kind='bar',
↳stacked=True);
```



Observation- female were more likely to survive more than men

```
[166]: ## Finding out the missing values

df.isna().sum().sort_values(ascending = False)
```

```
[166]: Cabin          687
Age             177
Embarked         2
Died              0
Fare              0
Ticket           0
Parch            0
SibSp            0
Sex              0
Name             0
Pclass           0
```

```
Survived      0
PassengerId   0
dtype: int64
```

```
[167]: # is age and cabin highly correlated with the label we are interested in
↳ predicting?
# checking correlation for that
df.corr()
```

```
[167]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	\
PassengerId	1.000000	-0.005007	-0.035144	0.036847	-0.057527	-0.001652	
Survived	-0.005007	1.000000	-0.338481	-0.077221	-0.035322	0.081629	
Pclass	-0.035144	-0.338481	1.000000	-0.369226	0.083081	0.018443	
Age	0.036847	-0.077221	-0.369226	1.000000	-0.308247	-0.189119	
SibSp	-0.057527	-0.035322	0.083081	-0.308247	1.000000	0.414838	
Parch	-0.001652	0.081629	0.018443	-0.189119	0.414838	1.000000	
Fare	0.012658	0.257307	-0.549500	0.096067	0.159651	0.216225	
Died	0.005007	-1.000000	0.338481	0.077221	0.035322	-0.081629	

	Fare	Died
PassengerId	0.012658	0.005007
Survived	0.257307	-1.000000
Pclass	-0.549500	0.338481
Age	0.096067	0.077221
SibSp	0.159651	0.035322
Parch	0.216225	-0.081629
Fare	1.000000	-0.257307
Died	-0.257307	1.000000

Age and cabin aren't really "correlated" with our label. And missing values aren't making them any more desirable. so I will drop Cabin, Age, Embarked.

```
[168]: # dropping Cabin, Age, Embarked from train data
df.drop({'Cabin', 'Age', 'Embarked'}, axis=1, inplace= True)

# dropping also from the test data set as well.

df_test.drop({'Cabin', 'Age', 'Embarked'}, axis=1, inplace= True)
```

```
[169]: #checking missing values in test
df_test.isna().sum().sort_values(ascending = False)
```

```
[169]: Fare      1
Ticket    0
Parch     0
SibSp     0
Sex       0
```

```
Name          0
Pclass        0
PassengerId   0
dtype: int64
```

```
[170]: # Setting up a loop to fill value for that specific row
```

```
for i in range(len(df_test['Fare'])):
    if df_test['PassengerId'][i] == 1044:
        df_test['Fare'][i] = 10
```

```
[171]: # Checking it if the value is filled up or not.
```

```
df_test.iloc[[152]]['Fare']
```

```
[171]: 152    10.0
      Name: Fare, dtype: float64
```

```
[172]: # checking updated shape
```

```
print(df.shape)
print(df_test.shape)
```

```
(891, 10)
```

```
(418, 8)
```

```
[173]: # checking first five entries of train data
```

```
df.head()
```

```
[173]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	SibSp	Parch	\
0	Braund, Mr. Owen Harris	male	1	0	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	1	0	
2	Heikkinen, Miss. Laina	female	0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	1	0	
4	Allen, Mr. William Henry	male	0	0	

	Ticket	Fare	Died
0	A/5 21171	7.2500	1
1	PC 17599	71.2833	0
2	STON/O2. 3101282	7.9250	0

3	113803	53.1000	0
4	373450	8.0500	1

```
[174]: # chekking entries of test data
df_test.head()
```

```
[174]:
```

	PassengerId	Pclass	Name	Sex	\
0	892	3	Kelly, Mr. James	male	
1	893	3	Wilkes, Mrs. James (Ellen Needs)	female	
2	894	2	Myles, Mr. Thomas Francis	male	
3	895	3	Wirz, Mr. Albert	male	
4	896	3	Hirvonen, Mrs. Alexander (Helga E Lindqvist)	female	

	SibSp	Parch	Ticket	Fare
0	0	0	330911	7.8292
1	1	0	363272	7.0000
2	0	0	240276	9.6875
3	0	0	315154	8.6625
4	1	1	3101298	12.2875

PassengerId later needed during predicting survuval on testing data, so saving it in another variable.

I need to remove Died from training data and PassengerId from both of the data sets that are less correlated with labels.

```
[175]: # saving PassengerId
id= df_test['PassengerId']
# Deleting PassengerId and Died
df.drop({'PassengerId', 'Died'}, axis=1, inplace= True)
df_test.drop({'PassengerId'}, axis=1, inplace= True)
```

```
[176]: #checking updated shape
print(df.shape)
print(df_test.shape)
```

(891, 8)

(418, 7)

The one extra column in train is the label. so lets take that out:

train_test_split the train then to get out validation data;

```
[177]: df.corr()
```

```
[177]:
```

	Survived	Pclass	SibSp	Parch	Fare
Survived	1.000000	-0.338481	-0.035322	0.081629	0.257307
Pclass	-0.338481	1.000000	0.083081	0.018443	-0.549500
SibSp	-0.035322	0.083081	1.000000	0.414838	0.159651
Parch	0.081629	0.018443	0.414838	1.000000	0.216225

```
Fare      0.257307 -0.549500  0.159651  0.216225  1.000000
```

data dimenisons are reduced so negative correlation is also reduced.

```
[178]: # Defining labels:
y= df['Survived']
# dropping it from trian:
df.drop({'Survived'}, axis= 1, inplace= True)
```

Since we have categorical variables, we must convert them to some sort of numeric value so that our model could understand and create a relationship between various attributes.

```
[179]: # converison of categorical to numerical:

df1= df
df2= df_test

df= pd.get_dummies(df)
df_test= pd.get_dummies(df_test)
```

```
[180]: for col in df.columns:
        if col not in df_test.columns:
            df.drop({col}, axis= 1, inplace= True)

for col in df_test.columns:
    if col not in df.columns:
        df_test.drop({col}, axis= 1, inplace= True)
```

```
[181]: # Checking out the shapes of both data sets:

print(df.shape)
print(df_test.shape)
```

```
(891, 123)
```

```
(418, 123)
```

```
[182]: from sklearn.model_selection import train_test_split

# Splitting data for training, validation

X_train, X_test, y_train, y_test= train_test_split(df, y, random_state= 42)
```

```
[183]: # Plotting the feature importances using the Boosted Gradient Descent
from xgboost import XGBClassifier
from xgboost import plot_importance

# Training the model
model = XGBClassifier()
```


```
model_importance = model.fit(X_train, y_train)
```

```
# Plotting the Feature importance bar graph
```

```
plt.rcParams['figure.figsize'] = [14,12]
```

```
sns.set(style = 'darkgrid')
```

```
plot_importance(model_importance);
```



Observation :

One who is paying high fair is more likely to survived.

Females are more likely to survive.

Data is ready for further classification after preprocessing. As I am here in this case focused on survival rate so I had done prprocessing of data in that relative manner. If we want to predict anything else, then we need to perform EDA (Exploratrty Data Analysis) to understand how to preprocess.

5 Clustering Text

```
[184]: import pandas as pd
import numpy as np
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
import re
from nltk.corpus import stopwords
import warnings
```

Neural Information Processing Systems (NIPS) is one of the top machine learning conferences in the world. It covers topics ranging from deep learning and computer vision to cognitive science and reinforcement learning. Original dataset available- <https://github.com/benhamner/nips-papers>

The code to scrape and create this dataset is on GitHub.

Dataset used in this assignment by extracted the paper text from the raw PDF files and are releasing that in CSV files.

```
[185]: #Read datasets/papers.csv into papers
papers = pd.read_csv("dataNips/papers.csv")

# Print out the first rows of papers
papers.head()
```

```
[185]:
```

	id	year	title	event_type	\
0	1	1987	Self-Organization of Associative Database and ...	NaN	
1	10	1987	A Mean Field Theory of Layer IV of Visual Cort...	NaN	
2	100	1988	Storing Covariance by the Associative Long-Ter...	NaN	
3	1000	1994	Bayesian Query Construction for Neural Network...	NaN	
4	1001	1994	Neural Network Ensembles, Cross Validation, an...	NaN	

	pdf_name	abstract	\
0	1-self-organization-of-associative-database-an...	Abstract Missing	
1	10-a-mean-field-theory-of-layer-iv-of-visual-c...	Abstract Missing	
2	100-storing-covariance-by-the-associative-long...	Abstract Missing	
3	1000-bayesian-query-construction-for-neural-ne...	Abstract Missing	
4	1001-neural-network-ensembles-cross-validation...	Abstract Missing	

	paper_text
0	767\n\nSELF-ORGANIZATION OF ASSOCIATIVE DATABA...
1	683\n\nA MEAN FIELD THEORY OF LAYER IV OF VISU...
2	394\n\nSTORING COVARIANCE BY THE ASSOCIATIVE\nn...
3	Bayesian Query Construction for Neural\nNetwor...
4	Neural Network Ensembles, Cross\nValidation, a...

```
[186]: print(papers.paper_text[0][:500] + ' ...')
```


SELF-ORGANIZATION OF ASSOCIATIVE DATABASE

AND ITS APPLICATIONS

Hisashi Suzuki and Suguru Arimoto

Osaka University, Toyonaka, Osaka 560, Japan

ABSTRACT

An efficient method of self-organizing associative databases is proposed together with applications to robot eyesight systems. The proposed databases can associate any input with some output. In the first half part of discussion, an algorithm of self-organization is proposed. From an aspect of hardware, it produces a new style of neural netwo
...

An example of a text stored in data file

```
[187]: papers = papers.drop(["id", "event_type", "pdf_name"], axis = 1)
papers.head()
```

```
[187]:   year      title      abstract \
0  1987  Self-Organization of Associative Database and ...  Abstract Missing
1  1987  A Mean Field Theory of Layer IV of Visual Cort...  Abstract Missing
2  1988  Storing Covariance by the Associative Long-Ter...  Abstract Missing
3  1994  Bayesian Query Construction for Neural Network...  Abstract Missing
4  1994  Neural Network Ensembles, Cross Validation, an...  Abstract Missing


      paper_text
0  767\n\nSELF-ORGANIZATION OF ASSOCIATIVE DATABA...
1  683\n\nA MEAN FIELD THEORY OF LAYER IV OF VISU...
2  394\n\nSTORING COVARIANCE BY THE ASSOCIATIVE\nn...
3  Bayesian Query Construction for Neural\nNetwor...
4  Neural Network Ensembles, Cross\nValidation, a...
```

For the analysis of the papers, I am only interested in the text data associated with the paper as well as the year the paper was published in.

I will analyse this text data using natural language processing. Since the file contains some meta-data, such as Id and file names, it is necessary to remove all columns that do not contain useful text information.

```
[188]: groups = papers.groupby('year')
counts = groups.size()

plt.figure(figsize = (13, 8))
ax = sns.barplot(counts.index, counts.values)
ax.set_xlabel("Year", fontsize = 15, labelpad = 15)
plt.xticks(rotation = 90)
plt.show()
```



This graph shows number of papers per year.

```
[189]: display(papers['title'].head())

papers['title_processed'] = papers['title'].map(lambda x: re.sub('[,\.!?', ' ',
↪x))
papers['title_processed'] = papers['title_processed'].map(str.lower)

display(papers['title_processed'].head())
```

```
0    Self-Organization of Associative Database and ...
1    A Mean Field Theory of Layer IV of Visual Cort...
2    Storing Covariance by the Associative Long-Ter...
3    Bayesian Query Construction for Neural Network...
4    Neural Network Ensembles, Cross Validation, an...
Name: title, dtype: object
```

```
0    self-organization of associative database and ...
1    a mean field theory of layer iv of visual cort...
2    storing covariance by the associative long-ter...
3    bayesian query construction for neural network...
4    neural network ensembles cross validation and ...
Name: title_processed, dtype: object
```

I used a regular expression to remove any punctuation in the title. Then I will perform lowercasing. I'll then print the titles of the first rows before and after applying the modification.

```
[190]: # plotting 10 most common numbers

from sklearn.feature_extraction.text import CountVectorizer

def plot_10_most_common_words(count_data, count_vectorizer):
    words = count_vectorizer.get_feature_names()
    total_counts = np.zeros(len(words))
    for t in count_data:
        total_counts += t.toarray()[0]

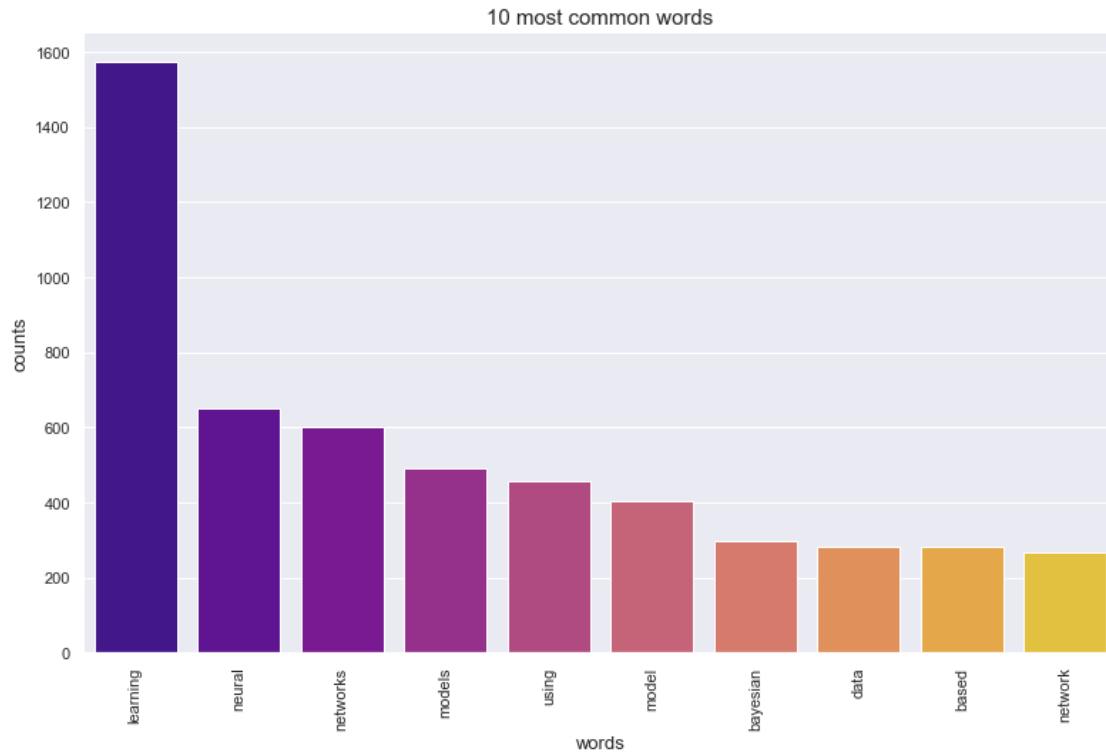
    count_dict = (zip(words, total_counts))
    count_dict = sorted(count_dict, key=lambda x:x[1], reverse=True)
    words = [w[0] for w in count_dict[0:10]]
    counts = [w[1] for w in count_dict[0:10]]
    x_pos = np.arange(len(words))

    sns.barplot(x_pos, counts, palette="plasma")
    plt.xticks(x_pos, words, rotation = 90)
    plt.xlabel('words', fontsize = 13)
    plt.ylabel('counts', fontsize = 13)
    plt.title('10 most common words', fontsize = 15)
    plt.show()

    return dict(count_dict)

count_vectorizer = CountVectorizer(stop_words = 'english')
count_data = count_vectorizer.fit_transform(papers['title_processed'])

plt.figure(figsize = (13, 8))
count_dict = plot_10_most_common_words(count_data, count_vectorizer)
```



```
[191]: #checking null values if exists.
papers.isnull().sum()
```

```
[191]: year          0
title             0
abstract          0
paper_text        0
title_processed   0
dtype: int64
```

```
[192]: # Cleaning unwanted
#Remove numbers,extra chracters, white spaces if they are not relevant to
↳ your analyses.
#Usually, regular expressions are used to remove numbers.

def cleaned_text(text):
    clean = re.sub("\n"," ",text)
    clean=clean.lower()
    clean=re.sub(r"[~.,%/:;?_&+*!=-]", " ",clean)
    clean=re.sub("[^a-z]", " ",clean)
    clean=clean.lstrip()
    clean=re.sub("\s{2,}", " ",clean)
```

```
    return clean
papers["cleaned_paper_text"]=papers["paper_text"].apply(cleaned_text)
```

```
[ ]:
```

```
[193]: # adding new column to exsisting dataframe as cleaned_paper_text
papers["cleaned_paper_text"] = papers["cleaned_paper_text"].apply(lambda x: ' '.
    ↳join([word for word in x.split() if len(word)>3]))
```

```
[194]: # display new column
papers["cleaned_paper_text"].head(10)
```

```
[194]: 0    self organization associative database applica...
      1    mean field theory layer visual cortex applicat...
      2    storing covariance associative long term poten...
      3    bayesian query construction neural network mod...
      4    neural network ensembles cross validation acti...
      5    sing neural instantiate deformable model chris...
      6    plasticity mediated competitive learning terre...
      7    iceg morphology classification using analogue ...
      8    real time control tokamak plasma using neural ...
      9    real time control tokamak plasma using neural ...
      Name: cleaned_paper_text, dtype: object
```

```
[ ]:
```

```
[195]: # Creating word cloud from cleaned data

from wordcloud import WordCloud
cloud=WordCloud(colormap="winter",width=600,height=400).
    ↳generate(str(papers["cleaned_paper_text"]))
fig=plt.figure(figsize=(13,18))
plt.axis("off")
plt.imshow(cloud,interpolation='bilinear')
```

```
[195]: <matplotlib.image.AxesImage at 0x2996f329f48>
```



```
[198]: papers.head()
```

```
[198]:   year          title          abstract \
0  1987  Self-Organization of Associative Database and ...  Abstract Missing
1  1987  A Mean Field Theory of Layer IV of Visual Cort...  Abstract Missing
2  1988  Storing Covariance by the Associative Long-Ter...  Abstract Missing
3  1994  Bayesian Query Construction for Neural Network...  Abstract Missing
4  1994  Neural Network Ensembles, Cross Validation, an...  Abstract Missing
```

```
          paper_text \
0  767\n\nSELF-ORGANIZATION OF ASSOCIATIVE DATABA...
1  683\n\nA MEAN FIELD THEORY OF LAYER IV OF VISU...
2  394\n\nSTORING COVARIANCE BY THE ASSOCIATIVE\n...
3  Bayesian Query Construction for Neural\nNetwor...
4  Neural Network Ensembles, Cross\nValidation, a...
```

```
          title_processed \
0  self-organization of associative database and ...
1  a mean field theory of layer iv of visual cort...
2  storing covariance by the associative long-ter...
3  bayesian query construction for neural network...
4  neural network ensembles cross validation and ...
```

```
          cleaned_paper_text \
0  self organization associative database applica...
1  mean field theory layer visual cortex applicat...
2  storing covariance associative long term poten...
3  bayesian query construction neural network mod...
4  neural network ensembles cross validation acti...
```

```
          stop_removed_paper_text
0  self organization associative database applica...
1  mean field theory layer visual cortex applicat...
2  storing covariance associative long term poten...
3  bayesian query construction neural network mod...
4  neural network ensembles cross validation acti...
```

```
[199]: papers['stop_removed_paper_text']
```

```
[199]: 0      self organization associative database applica...
      1      mean field theory layer visual cortex applicat...
      2      storing covariance associative long term poten...
      3      bayesian query construction neural network mod...
      4      neural network ensembles cross validation acti...

      ...
      7236     single transistor learning synapses paul hasle...
      7237     bias variance combination least squares estima...
```

```

7238    real time clustering cmos neural engine serran...
7239    learning direction global motion classes psych...
7240    correlation interpolation networks real time e...
Name: stop_removed_paper_text, Length: 7241, dtype: object

```

Tokenization Tokenization is the process of splitting the given text into smaller pieces called tokens. Words, numbers, punctuation marks, and others can be considered as tokens.

```
[200]: papers["tokenized"]=papers["stop_removed_paper_text"].apply(lambda x: nltk.
      ↪word_tokenize(x))
```

```
[201]: papers.head()
```

```
[201]:
```

	year		title	abstract	\
0	1987	Self-Organization of Associative Database and ...	Abstract Missing		
1	1987	A Mean Field Theory of Layer IV of Visual Cort...	Abstract Missing		
2	1988	Storing Covariance by the Associative Long-Ter...	Abstract Missing		
3	1994	Bayesian Query Construction for Neural Network...	Abstract Missing		
4	1994	Neural Network Ensembles, Cross Validation, an...	Abstract Missing		

```

                                paper_text \
0  767\n\nSELF-ORGANIZATION OF ASSOCIATIVE DATABA...
1  683\n\nA MEAN FIELD THEORY OF LAYER IV OF VISU...
2  394\n\nSTORING COVARIANCE BY THE ASSOCIATIVE\n...
3  Bayesian Query Construction for Neural\nNetwor...
4  Neural Network Ensembles, Cross\nValidation, a...

```

```

                                title_processed \
0  self-organization of associative database and ...
1  a mean field theory of layer iv of visual cort...
2  storing covariance by the associative long-ter...
3  bayesian query construction for neural network...
4  neural network ensembles cross validation and ...

```

```

                                cleaned_paper_text \
0  self organization associative database applica...
1  mean field theory layer visual cortex applicat...
2  storing covariance associative long term poten...
3  bayesian query construction neural network mod...
4  neural network ensembles cross validation acti...

```

```

                                stop_removed_paper_text \
0  self organization associative database applica...
1  mean field theory layer visual cortex applicat...
2  storing covariance associative long term poten...
3  bayesian query construction neural network mod...
4  neural network ensembles cross validation acti...

```



```

                                tokenized
0  [self, organization, associative, database, ap...
1  [mean, field, theory, layer, visual, cortex, a...
2  [storing, covariance, associative, long, term,...
3  [bayesian, query, construction, neural, networ...
4  [neural, network, ensembles, cross, validation...

```

We can see a column added in this dataset as tokenized.

Lematization The aim of lemmatization, like stemming, is to reduce inflectional forms to a common base form. As opposed to stemming, lemmatization does not simply chop off inflections. Instead it uses lexical knowledge bases to get the correct base forms of words.

```

[202]: # This step take some time ....In my pc it taken around 20 minutes.
import nltk
nltk.download('wordnet')

def word_lemmatizer(text):
    lem_text = [WordNetLemmatizer().lemmatize(i,pos='v') for i in text]
    return lem_text
papers["lemmatized"]=papers["tokenized"].apply(lambda x: word_lemmatizer(x))
papers["lemmatize_joined"]=papers["lemmatized"].apply(lambda x: ' '.join(x))

```

```

[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\DeLL\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!

```

```

[203]: papers["lemmatize_joined"].head()

```

```

[203]: 0    self organization associative database applica...
1    mean field theory layer visual cortex applicat...
2    store covariance associative long term potenti...
3    bayesian query construction neural network mod...
4    neural network ensembles cross validation acti...
Name: lemmatize_joined, dtype: object

```

```

[204]: # adding a new column for data analysis
papers['Number_of_words_for_cleaned'] = papers['lemmatize_joined'].apply(lambda_
↪x:len(str(x).split()))


```

```

[205]: # plotting 30 most frequent words after text preprocessing.
plt.style.use('ggplot')
plt.figure(figsize=(14,6))
freq=pd.Series(" ".join(papers["lemmatize_joined"]).split()).value_counts()[:30]
freq.plot(kind="bar", color = "orangered")
plt.title("30 most frequent words",size=20)

```

```
[205]: Text(0.5, 1.0, '30 most frequent words')
```



This is after preprocessing of text file. As before preprocessing it was different

After the text preprocessing is done, this result may be used for more complicated NLP tasks, for example, machine translation or natural language generation.

```
[ ]:
```

```
[ ]:
```